



Concurrency vs. Parallelism

□ Concurrent execution on single-core system:



□ Parallelism on a multi-core system:



Multicore Programming (Cont.)

- Types of parallelism
 - **Data parallelism** – distributes subsets of the data across multiple cores, same operation on each
 - **Task parallelism** – distributing threads across cores, each thread performing unique operation


Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- S is serial portion
- N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$

Serial portion of an application has disproportionate effect on performance gained by adding additional cores



User Threads and Kernel Threads


- **User threads** – Thread management takes place using a threads library without OS support.
 - Kernel would treat the process as single-threaded and any blocking system call by one of the threads would end up blocking all the threads of that process.
- **Kernel threads** - Supported and managed by the OS Kernel. Kernel support exists for most well-known Oses, e.g.:
 - Windows
 - Solaris
 - Linux
 - Tru64 UNIX
 - Mac OS X
- **NOTE:** A kernel thread is not meant to indicate a thread executing kernel code, but rather a thread that is managed and supported by the kernel.

4.3 Multithreading Models

- Several thread management models exist:
 - Many-to-One model
 - One-to-One model
 - Many-to-Many model


Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread **blocking** causes all to block
- Multiple threads may **not run in parallel** on multicore systems because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - **Solaris Green Threads**
 - **GNU Portable Threads**




One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux
 - Solaris 9 and later




Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads;
 $\# \text{ kernel threads} \leq \# \text{ user threads}$
- Allows the operating system to create a sufficient number of kernel threads
- Examples:
 - Windows with the *ThreadFiber* package
 - Solaris prior to version 9



Many-to-Many variant: The Two-level Model

- Similar to the many-to-many model in that it allows many user threads to map to many kernel threads.
- But it also allows a one-to-one relationship on some user/kernel thread pairs as shown on the diagram.
- Examples
 - HP-UX
 - Tru64 UNIX
 - Solaris prior to version 8



4.4 Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads and is responsible for implementing a thread management model
- Three primary thread libraries:
 - POSIX **Pthreads** (whether user/kernel thread is platform dependent, but same interface)
 - Windows threads (kernel threads)
 - Java threads (user threads)
- Two primary ways of implementing
 - Library entirely in user space (i.e. with no kernel support)
 - Kernel-level library supported by the OS

4.4.1 Pthreads

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- May be implemented either as user-level or kernel-level.
- ***Specification***, not ***implementation***
 - Different implementation for different OS platforms, but all have the same interface.
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX-like operating systems (Solaris, Linux, Mac OS X)

Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

```
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* create the thread */
    pthread_create(&tid,&attr,runner,argv[1]);
    /* wait for the thread to exit */
    pthread_join(tid,NULL);

    printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

4.4.2 Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr, "An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr, "An integer >= 0 is required\n");
        return -1;
    }
}
```

Windows Multithreaded C Program (Cont.)

```
/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */

if (ThreadHandle != NULL) {
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}
}
```


Windows Multithreaded C Program (Cont.)

The creating thread can use the arguments to [CreateThread](#) to specify the following:

- The security attributes for the handle to the new thread, including:
 - An **inheritance flag** that determines whether the handle can be inherited by child processes.
 - A **security descriptor**, which the system uses to perform access checks on all subsequent uses of the thread's handle before access is granted.
- The initial **stack size** of the new thread.
 - The thread's stack is allocated automatically in the memory space of the process
 - The system increases the stack as needed and frees it when the thread terminates.
- A creation flag that enables you to create the thread in a **suspended** state. When suspended, the thread does not run until the [ResumeThread](#) function is called.

4.5 Implicit Threading

- Growing in popularity; As the number of threads increases, program correctness becomes more difficult (with explicit threads).
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- Three methods explored:
 - Thread Pools
 - OpenMP
 - Grand Central Dispatch
- Other methods include Microsoft Threading Building Blocks (TBB) and **java.util.concurrent** package

4.5.1 Thread Pools

- **At process startup**, create a **predefined** number of threads in a pool where they await work.
- Threads wait for work to be dispatched to them. When work is dispatched to the thread it performs it and when done returns back to the pool.
- If more work needs to be dispatched with no threads available in the pool, the main process waits till a thread becomes available to the pool.
- Advantages:
 - Servicing a request with an existing thread is **faster than creating** a new thread
 - Allows the number of threads in the application(s) to be **bound** to the size of the pool (thus preventing the creation of too many threads)
 - Separating tasks to be performed **from mechanics of creating task allows** different strategies for running tasks
 - e.g. Tasks may be scheduled to run as **one-shot** after a delay time
 - or may be scheduled to run **periodically**
- Works very well for tasks that have a finite duration, i.e. tasks that start, do some work, then exit.

Thread Pools – cont.

- **The number of threads** in a pool may be **pre-determined** based on the number of CPUs, memory or the expected number of client requests.
- Alternatively, more sophisticated systems may adjust the number of threads **dynamically**.
- Windows API supports thread pools – The user may call an API for dispatching work to a thread using the function:

```
BOOL QueueUserWorkItem(LPTHREAD_START_ROUTINE  
    Function, PVOID Param, ULONG Flags );
```

- The function may take the form:

```
DWORD WINAPI PoolFunction(PVOID Param) {  
    /* this function runs as a separate thread */  
}
```

4.5.2 OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel using:
#pragma omp parallel
- Creates as many threads as there are cores

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```

OpenMP cont.

- The following runs a for-loop in parallel

```
#pragma omp parallel
for (i=0; i<N; i++) {
    c[i] = a[i] + b[i];
}
```

- OpenMP allows developers to control the level of parallelism by allowing:
 - Manual setting of the **number of threads**.
 - Specify certain data as **shared** or private
- OpenMP is available for open-source as well as commercial compilers:
 - Linux, Windows and Mac OS X.

4.6 Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations

4.6.1 Semantics of `fork()` and `exec()`

- Does **`fork()`** duplicate only the calling thread or all threads?
 - Some UNIX systems have two versions of `fork`
 - In Linux, if a process has multiple threads and one of them calls `fork()`, the child thread will have a replica of the parent's code, data, stack, heap, file and other resources BUT will only have one thread running, the one that called the `fork()` function.
- **`exec()`** usually works as normal – replace the running process **including all its threads**

4.6.2 Signal Handling

- **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- Synchronous signals are those caused by the running process (e.g. divide by zero, or memory illegal memory access). Asynchronous signals are caused outside the program (e.g. upon an expiration of a timer after the process issues an `alarm()` call).
- A **signal handler** is a function that is used to process/handle signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Each signal is handled by one of two signal handlers:
 1. default
 2. user-defined (provided by the process)
- Every signal has **default handler** that kernel runs when handling signal
 - **User-defined signal handler** can override default
 - For single-threaded, the signal is delivered to process's main thread → no issues

Signal Handling (Cont.)

- Where should a signal be delivered for multi-threaded?
 - Deliver the signal to the thread to which the signal **applies**
 - Deliver the signal to **every** thread in the process
 - Deliver the signal to **certain threads** in the process
 - Assign a **specific thread** to receive **all** signals for the process
- The standard linux function for delivering signals is:
`Kill (pid_t pid, int signal)`
- Most Unix versions allow a **thread to specify which signals it accepts** (else the signal is blocked), and a signal is **delivered ONLY ONCE to the first** thread that accepts it.
- POSIX allows signal delivery to a specified thread
`pthread_kill(pthread_t tid, int signal);`
- Windows (which doesn't have signals) uses Asynchronous procedure calls (APC)
 - Allows a thread X to specify an APC function to another thread Y.

```
DWORD QueueUserAPC(PAPCFUNC pfnAPC, HANDLE hThread, ULONG_PTR dwData );
```

- Windows calls the APC function once thread Y blocks for an event, semaphore or any synchronization object.

4.6.3 Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled and thus terminates orderly.
- Difficulty: When a thread is allocated resources or shared data. When canceling a thread asynchronously, the OS may reclaim only system resources. Also a thread may be in the middle of updating some shared data.

Thread Cancellation - cont.

- Pthread code to create and cancel a thread:

```
pthread_t tid;  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
...  
/* cancel the thread */  
pthread_cancel(tid);
```

Thread Cancellation (Cont.)

- In pThreads, invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled, cancellation remains pending until thread enables it.
 - `pthread_setcanceltype(..)` – to set thread's mode
 - `pthread_setcancelstate(..)` – to set thread's state
- Default is deferred and enabled
 - Cancellation only occurs when thread reaches **cancellation point**,
 - By calling `pthread_test_cancel()` to establish a cancellation point (after it frees the assigned resources or gracefully stops manipulating shared data).
 - After calling `pthread_join()`
 - After calling `sigwait()` or `pthread_cond_wait()`
- On Linux systems, thread cancellation is handled through signals

4.6.4 Thread-Local Storage

- Generally threads within the same process share their data, but sometimes a thread may have some data that is not to be shared (e.g. tabbed webpages each with a thread) -> it needs **Thread-local storage (TLS)**.
- Local variables within a function do not directly achieve that, since they do not propagate across function calls.
- Most thread libraries (pthreads and Windows API) provide some form of support.
- For pthreads use the `__thread` specifier prior to variable declaration, e.g.

```
__thread int i;
```

This requires significant support from the compiler (gcc), linker (ld), dynamic linker (ld.so) and system libraries (libc.so and libpthread.so)


Race Conditions

e.g. if #counter = 0xF230

```
int counter=0;
void thread1()
{
    .
    .

    counter++;

    .
    .
}
```



```
void thread2()
{
    .
    .

    counter--;

    .
    .
}
```

Memory

Race Conditions

counter++ could be implemented in machine code as

```
mov r2, #counter
```

```
ld r1, [r2]
```

```
inc r1
```

```
st r1, [r2]
```

```
register1 = counter
```

```
register1 = register1 + 1
```

```
counter = register1
```

• **counter--** could be implemented in machine code as

```
mov r3, #counter
```

```
ld r2, [r3]
```

```
dec r2
```

```
st r2, [r3]
```

```
register2 = counter
```

```
register2 = register2 - 1
```

```
counter = register2
```

• Consider this execution interleaving with “count = 5” initially:

S0: producer execute **register1 = counter**

{register1 = 5}

S1: producer execute **register1 = register1 + 1**

{register1 = 6}

S2: consumer execute **register2 = counter**

{register2 = 5}

S3: consumer execute **register2 = register2 - 1**

{register2 = 4}


S4: producer execute **counter = register1**

{counter = 6}

S5: consumer execute **counter = register2**

{counter = 4}

e.g. if #counter = 0xF230



Race Conditions – cont.

- A **race condition** or **race hazard** is the behavior of a software (or hardware) system where the output is dependent on the sequence or relative timing of the executing threads.
- A **critical race condition** occurs when the order of operations on shared variables causes them to have unexpected or **erroneous** values.
- A **non-critical race condition** occurs when the order of operations on shared variables does not result in an unexpected or erroneous value.
- Critical race conditions result in invalid execution and bugs. Failure to obey mutual exclusion opens up the possibility of corrupting the shared variables.

Race Conditions - cont.

- A critical race condition occurs when **multiple threads** are performing **non-atomic read-modify-write** concurrently or in parallel.
- It is not necessary for two threads writing to a shared variable concurrently, to result in a critical race condition. A read-modify-write needs to exist to cause a critical race condition.
- Examples of read-modify-write operations:
 - Increment and decrement (e.g. `counter++`, `counter--`)
 - test-and-set
 - compare-and-swap
 - accumulate operations (e.g. `counter+=4`)

Race Conditions – cont.

- Race conditions have a reputation of being [difficult to reproduce and debug](#), since the end result is nondeterministic and depends on the relative timing between interfering threads.
- Problems occurring in production systems can therefore disappear when running in debug mode, when additional logging is added, or when attaching a debugger. Thus, a bug that is due to a race condition is often referred to as a ["Heisenbug"](#).
- Thus, it is better to avoid race conditions in the first place and there is no alternative to proper and careful software design.

5.2 Critical Section Problem

- Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- Each process has **critical section** segment of code (the section that manipulates shared variables using read-modify-write operations)
 - A Process may be changing common variables, updating a table, writing file, etc
 - To avoid race conditions, when one process is in critical section, no other should be in its critical section
- **Critical section problem** is to design protocol to solve this
- Each process must ask permission to enter critical section. This happens in the **entry section**. It may follow critical section with an **exit section**, then **remainder section**

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then:
 - The selection of the processes that will enter the critical section next cannot be postponed indefinitely.
 - Only processes that are in their entry section can participate in the selection.
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section (and before its request is granted)
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes

5.3 Peterson's Solution

- Good algorithmic description of solving the problem
- **Two process solution** (i.e. works for two processes only)
- Assume that the `load` and `store` machine-language instructions are atomic; that is, cannot be interrupted (a reasonable assumption)
- The two processes share two variables:
 - `int turn;`
 - `bool req[2]`
- The variable `turn` indicates whose turn it is to enter the critical section
- The `req` array is used to indicate if a process is ready to enter the critical section. `req[i] = true` implies that process P_i is ready!

Algorithm for Process P_i

```
do {  
    req[i] = true;  
    turn = j;  
    while (req[j] && turn == j);  
    critical section  
    req[i] = false;  
    remainder section  
} while (true);
```

Algorithm for Process P_i

```
do {  
    req[0] = true;  
    turn = 1;  
    while (req[1] && turn == 1);  
    critical section  
    req[0] = false;  
    remainder section  
} while (true);
```

```
do {  
    req[1] = true;  
    turn = 0;  
    while (req[0] && turn == 0);  
    critical section  
    req[1] = false;  
    remainder section  
} while (true);
```


Peterson's Solution (Cont.)

- Provable that the three CS requirement are met:
 1. Mutual exclusion is preserved

P_i enters CS only if:
either **reg[j]==false** or **turn==i**
 2. Progress requirement is satisfied
 3. Bounded-waiting requirement is met
- **Note:** The two threads are setting the turn variable. This is **not** a read-modify-write and does not result in a critical race condition.

5.4 Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- All H/W solutions described in this section are based on idea of **locking**
 - Protecting critical regions via locks
- **Uniprocessors** – could **disable interrupts**
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems since it requires sending a disable interrupts message to all cores.
 - Operating systems using this approach are not broadly scalable
- Modern machines provide special **atomic hardware instructions**
 - **Atomic** = non-interruptible
 - Either test memory word and set value
 - Or swap contents of two memory words

Solution to Critical-section Problem Using Locks

Process A

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder  
section  
} while (TRUE);
```

Process B

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder  
section  
} while (TRUE);
```

test_and_set Instruction

Definition:

```
bool test_and_set (bool *target)
{
    bool rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically (**it is a single machine instruction**)
it is a single machine instruction
1. Returns the original value of the lock variable (*target)
2. Set the new value of lock variable (*target) to “TRUE”.

Using test_and_set()

- Shared Boolean variable lock, initialized to FALSE
- A possible solution to critical section problem?

```
do {  
    /* Wait till lock is false i.e. not locked, then acquire it */  
    while (test_and_set(&lock));  
  
    /* critical section */  
    . . .  
    /* release the lock at the end (i.e. make it false) */  
    lock = false;  
  
    /* remainder section */  
    . . .  
  
} while (true);
```

fetch_and_add Instruction

Definition:

```
int fetch_and_add (int *target, int inc)
{
    int rv = *target;
    *target = *target + inc;
    return rv;
}
```

1. Executed atomically (**it is a single machine instruction**)
2. Returns the original value of the lock variable (*target)
3. Set the new value of (*target) to (*target) + inc.

compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int rv = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return rv;  
}
```

1. Executed atomically
2. Returns the original value of the lock variable (`*value`)
3. Set the variable “value” the value of the passed parameter “new_value” but only if “*value” == “expected”. That is, the swap takes place only under this condition.

Using compare_and_swap

- Shared integer “lock” initialized to 0;
- A possible solution to critical section problem?

```
do {  
    /* Wait for value to be zero (i.e. lock is released), then acquire lock */  
    while (compare_and_swap(&lock, 0, 1) != 0);  
  
    /* critical section */  
    . . .  
    /* release the lock when done with CS */  
    lock = 0;  
  
    /* remainder section */  
    . . .  
} while (true);
```


Bounded-waiting Mutual Exclusion with test_and_set

- **Previous algorithms didn't satisfy the bounded wait requirement.**
- This algorithm uses common data structures:

```
bool waiting[n];  
bool lock;
```
- The variable `Key` is not shared
- Proof of mutual exclusion:
 - P_i can enter its critical section only if either `waiting[i] == false` OR `key == false`.
 - The value of `key` can become false only if `test_and_set()` is executed. The first process to execute it will find `key == false`; all others must wait.
 - The variable `waiting[i]` can become false only if another process leaves its critical section; only one `waiting[i]` is set to false, maintaining the mutual-exclusion requirement.

```
do {
```

```
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;
```

```
    /* critical section */
```

```
    ...
```

```
    /* Select next process to run  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
  
    if (j == i)  
        lock = false;  
    else  
        waiting[j] = false;
```

```
    /* remainder section starts below*/
```

```
    ...
```

```
    } while (true);
```

Bounded-waiting Mutual Exclusion with test_and_set

- Proof of progress:
 - Since a process exiting the critical section either sets `lock` to `false` or sets `waiting[j]` to `false`. Both allow a process that is waiting to enter its critical section to proceed.
- Proof of bounded wait:
 - When a process leaves its critical section, it scans the array `waiting` in the cyclic ordering $(i + 1, i + 2, \dots, n - 1, 0, \dots, i - 1)$. It designates the first process in this ordering that is in the entry section (`waiting[j] == true`) as the next one to enter the critical section. Any process waiting to enter its critical section will thus do so within $n - 1$ turns.

```
do {
```

```
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
```

```
    /* critical section */
```

```
    ...
```

```
    /* Select next process to run
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;

    if (j == i)
        lock = false;
    else
        waiting[j] = false;
```

```
    /* remainder section starts below*/
```

```
    ...
```

```
    } while (true);
```