

5.3 Peterson's Solution

- Good algorithmic description of solving the problem
- **Two process solution** (i.e. works for two processes only)
- Assume that the `load` and `store` machine-language instructions are atomic; that is, cannot be interrupted (a reasonable assumption)
- The two processes share two variables:
 - `int turn;`
 - `bool req[2]`
- The variable `turn` indicates whose turn it is to enter the critical section
- The `req` array is used to indicate if a process is ready to enter the critical section. `req[i] = true` implies that process P_i is ready!

Algorithm for Process P_i

```
do {  
    req[i] = true;  
    turn = j;  
    while (req[j] && turn == j);  
    critical section  
    req[i] = false;  
    remainder section  
} while (true);
```

Algorithm for Process P_i

```
do {  
    req[0] = true;  
    turn = 1;  
    while (req[1] && turn == 1);  
    critical section  
    req[0] = false;  
    remainder section  
} while (true);
```

```
do {  
    req[1] = true;  
    turn = 0;  
    while (req[0] && turn == 0);  
    critical section  
    req[1] = false;  
    remainder section  
} while (true);
```

Peterson's Solution (Cont.)

- Provable that the three CS requirement are met:
 1. Mutual exclusion is preserved
 - P_i enters CS only if:
either **`reg[j]==false`** or **`turn==i`**
 2. Progress requirement is satisfied
 3. Bounded-waiting requirement is met
- **Note:** The two threads are setting the turn variable. This is **not** a read-modify-write and does not result in a critical race condition.

5.4 Synchronization Hardware

- Many systems provide hardware support for implementing the critical section code.
- All H/W solutions described in this section are based on idea of **locking**
 - Protecting critical regions via locks
- **Uniprocessors** – could **disable interrupts**
 - Currently running code would execute without preemption
 - Generally too inefficient on multiprocessor systems since it requires sending a disable interrupts message to all cores.
 - Operating systems using this approach are not broadly scalable
- Modern machines provide special **atomic hardware instructions**
 - **Atomic** = non-interruptible
 - Either test memory word and set value
 - Or swap contents of two memory words

Solution to Critical-section Problem Using Locks

Process A

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder  
section  
} while (TRUE);
```

Process B

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder  
section  
} while (TRUE);
```

test_and_set Instruction

Definition:

```
bool test_and_set (bool *target)
{
    bool rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically (**it is a single machine instruction**)
it is a single machine instruction
1. Returns the original value of the lock variable (*target)
2. Set the new value of lock variable (*target) to “TRUE”.

Using test_and_set()

- Shared Boolean variable lock, initialized to FALSE
- A possible solution to critical section problem?

```
do {  
    /* Wait till lock is false i.e. not locked, then acquire it */  
    while (test_and_set(&lock));  
  
    /* critical section */  
    . . .  
    /* release the lock at the end (i.e. make it false) */  
    lock = false;  
  
    /* remainder section */  
    . . .  
  
} while (true);
```


fetch_and_add Instruction

Definition:

```
int fetch_and_add (int *target, int inc)
{
    int rv = *target;
    *target = *target + inc;
    return rv;
}
```

1. Executed atomically (**it is a single machine instruction**)
2. Returns the original value of the lock variable (`*target`)
3. Set the new value of (`*target`) to `(*target) + inc`.

compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int rv = *value;  
  
    if (*value == expected)  
        *value = new_value;  
    return rv;  
}
```

1. Executed atomically
2. Returns the original value of the lock variable (`*value`)
3. Set the variable “value” the value of the passed parameter “new_value” but only if “*value” == “expected”. That is, the swap takes place only under this condition.

Using compare_and_swap

- Shared integer “lock” initialized to 0;
- A possible solution to critical section problem?

```
do {  
    /* Wait for value to be zero (i.e. lock is released), then acquire lock */  
    while (compare_and_swap(&lock, 0, 1) != 0);  
  
    /* critical section */  
    . . .  
    /* release the lock when done with CS */  
    lock = 0;  
  
    /* remainder section */  
    . . .  
} while (true);
```

Bounded-waiting Mutual Exclusion with test_and_set

- **Previous H/W algorithms didn't satisfy the bounded wait requirement.**
- This algorithm uses common data structures:

```
bool waiting[n];  
bool lock;
```
- The variable `Key` is not shared
- Proof of mutual exclusion:
 - P_i can enter its critical section only if either `waiting[i] == false` OR `key == false`.
 - The value of `key` can become false only if `test_and_set()` is executed. The first process to execute it will find `key == false`; all others must wait.
 - The variable `waiting[i]` can become false only if another process leaves its critical section; only one `waiting[i]` is set to false, maintaining the mutual-exclusion requirement.

```
do {
```

```
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;
```

```
    /* critical section */
```

```
    ...
```

```
    /* Select next process to run  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
  
    if (j == i)  
        lock = false;  
    else  
        waiting[j] = false;
```

```
    /* remainder section starts below*/
```

```
    ...
```

```
    } while (true);
```

Bounded-waiting Mutual Exclusion with test_and_set

- Proof of progress:
 - Since a process exiting the critical section either sets `lock` to false or sets `waiting[j]` to false. Both allow a process that is waiting to enter its critical section to proceed.
- Proof of bounded wait:
 - When a process leaves its critical section, it scans the array `waiting` in the cyclic ordering $(i + 1, i + 2, \dots, n - 1, 0, \dots, i - 1)$. It designates the first process in this ordering that is in the entry section (`waiting[j] == true`) as the next one to enter the critical section. Any process waiting to enter its critical section will thus do so within $n - 1$ turns.

```
do {
```

```
    waiting[i] = true;  
    key = true;  
    while (waiting[i] && key)  
        key = test_and_set(&lock);  
    waiting[i] = false;
```

entry section



```
    /* critical section */
```

```
    ...
```

```
    /* Select next process to run  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
  
    if (j == i)  
        lock = false;  
    else  
        waiting[j] = false;
```

exit section



```
    /* remainder section starts below*/
```

```
    ...
```

```
    } while (true);
```

5.5 Mutex Locks

- The OS provides abstraction for the hardware tools previously described, particularly since they require some shared lock variables.
- Simplest is **mutex**.
- Usage: Protect a critical section by first **acquire()** a lock then **release()** the lock
 - Lock = Boolean variable indicating if lock is available or not

```
int main() {  
    do {  
        acquire lock  
        critical section  
        release lock  
        remainder section  
    } while (true);  
}
```

An implementation using **atomic acquire()** and **release()**

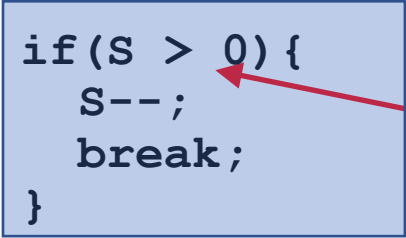
- May be implemented via hardware atomic instructions such as
 - `test_and_set` or `compare_and_swap`
- This lock sometimes referred to as a **spinlock** because it requires **busy waiting**, thus
 - **NOT EFFICIENT.**
 - When used, the critical section must be very short
- A mutex may also be implemented without a spinlock by using **wait queues**. The method is explained in the next section for semaphores.


```
int main() {  
    do {  
        acquire lock  
        critical section  
        release lock  
        remainder section  
    } while (true);  
}
```

5.6 Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore S – integer variable
- **Theoretically**, it can only be accessed via **indivisible (atomic) operations (shown in blue rectangles) `wait()` and `signal()`** (Originally called $P()$ and $V()$)

These are NOT UNIX `wait()` and `signal()` API calls

```
wait(S) {  
    while(true) { // busy wait till S>0  
          
    }  
}
```

```
signal(S) {  
      
}
```

$S > 0$ (i.e. 1 and above) indicates that the semaphore is not locked

BLUE RECTANGLES indicate atomic operations

Semaphore Usage

- **Counting semaphore usage** – integer value can range over an unrestricted domain
 - May be used to organize usage of a resource that only allows access to N processes at a time \rightarrow semaphore needs to be initialized to N .
- **Binary semaphore usage** – integer value can range only between 0 and 1
 - Same as a **mutex lock**, except that it has a different polarity (**initialized to 1**)
 - Can synchronization two processes: Consider two processes P_1 and P_2 that require a statement S_1 to happen before S_2

Create a semaphore named “**synch**” and **initialize it to 0**

P1 :

S_1 ;

signal (synch) ;

P2 :

wait (synch) ;

S_2 ;

- **Note that** generally you cannot initialize the semaphore’s value to less than zero. See the man page for `sem_init()`.


Semaphore Implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore **concurrently** (i.e. blue rectangles must be guaranteed to execute **atomically**)
- In similarity to a mutex, processes and threads can be **busy waiting** for the semaphore to become available (i.e. >0)
- A process may thus spend a lot of time in entry sections waiting for the semaphore and not doing any useful work.
- Therefore it may be efficient from the system's point of view to block the process and move it into a waiting queue, and schedule another ready process to run instead.

Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each semaphore has two data items:
 - value (of type integer)
 - pointer to first process in the linked-list queue.
- We define two operations:
 - **block** – place the process invoking the operation on the appropriate waiting queue
 - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```



Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {
```

```
    S->value--;
```

```
    if (S->value < 0) {
```

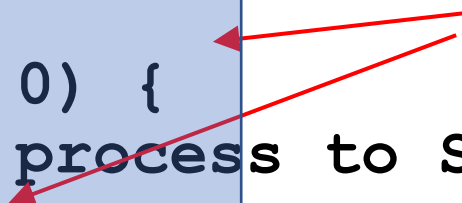
```
        /* add this process to S->list; */
```

```
        block();
```

```
    }
```

```
}
```

Reverse order compared to
that used in busy-waiting



BLUE RECTANGLES indicate
atomic operations

```
signal(semaphore *S) {
```

```
    S->value++;
```

```
    if (S->value <= 0) { /* if someone is in wait queue */
```

```
        /* remove a process P from S->list; */
```

```
        wakeup(P);
```

```
    }
```

```
}
```

Implementation with no Busy waiting (Cont.)

In this implementation, semaphore values are:

- <0 : indicates one or more processes are blocked waiting on the semaphore
 - This is different from previous implementation where the value cannot be <0 .
- $=0$: indicates the semaphore is not available but no process is blocked waiting on it.
- >0 : (i.e. 1 and above) indicates the semaphore is available and thus no process is blocked waiting on it.

Unix/Pthreads Synchronization

- Named semaphores use names that start with '/' and are less than 251 characters.

```
sem_open  
sem_post  
sem_wait  
sem_close  
sem_unlink
```

- Unnamed (anonymous) semaphores use shared variables (for processes or threads) of type `sem_t`.

```
sem_init  
sem_post  
sem_wait  
sem_destroy
```

Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let S and Q be two semaphores initialized to 1

P_0
`wait(S) ;`
`wait(Q) ;`
`...`
`signal(S) ;`
`signal(Q) ;`

P_1
`wait(Q) ;`
`wait(S) ;`
`...`
`signal(Q) ;`
`signal(S) ;`

Priority inversion

- **Priority inversion** is a Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - Consider having 3 processes L,M and H with low, medium and high priorities respectively, and consider also a resource R that is shared amongst them.
 - If process L acquired R, and then process H requested R, then H will be blocked.
 - If another process M (priority higher than L and is not requesting R) is ready to run, then it may preempt process L (due to the timer tick for example).
 - This indirectly causes priority inversion and it is sometimes problematic.
- Solved via **priority-inheritance protocol**
 - Priority of process L changes to high when H requests the shared resource, and thus won't be preempted by process M.
- This problem occurred on the Mars Pathfinder's robot in 1997 (running a VxWorks real-time OS).

5.7 Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
 - Bounded-Buffer Problem
 - Readers and Writers Problem
 - Dining-Philosophers Problem

Bounded-Buffer Problem

- ***BUF_SZ*** elements inside the shared buffer
- Semaphore **`num_full_el`** initialized to the value 0 – keeps track of the number of elements that are full.
- Semaphore **`num_empty_el`** initialized to the value n – keeps track of number of elements that are empty.
- Why use semaphores, when the previous approach seemed to work fine (i.e. using the in and out indices without using any synchronization primitives)?

Bounded-Buffer Problem

- **BUF_SZ** elements inside the shared buffer
- Semaphore **num_full_el** initialized to the value 0 – keeps track of the number of elements that are full.
- Semaphore **num_empty_el** initialized to the value n – keeps track of number of elements that are empty.
- Why use semaphores, when the previous approach seemed to work fine (i.e. using the in and out indices without using any synchronization primitives)?
 - Because of the busy-waiting problem in which a process or thread may be spending valuable CPU time doing nothing but waiting in a loop!
 - We may still use the in and out variables to index a particular buffer in the pool.
- The full semaphore is used to indicate how many buffers are full, whereas the empty semaphore indicates how many are empty.

Bounded Buffer Problem (Cont.)

- The structure of the **producer** process

```
do {  
    /* produce an item in next_produced */  
    ...  
    /* dec empty sem. */  
    wait(num_empty_el);  
    /* write/produce to an entry*/  
    buffer[in] = next_produced;  
    in = (in + 1) % BUF_SZ;  
    /* inc full sem. */  
    signal(num_full_el);  
} while (true);
```

- The structure of the **consumer** process

```
do {  
    /* dec full sem. */  
    wait(num_full_el);  
    /* read/consume an entry */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUF_SZ;  
    /*inc empty sem.*/  
    signal(num_empty_el);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```

Readers-Writers Problem

- A data set (e.g. a database) is shared among a number of concurrent processes
 - Readers – only read the data set; they do **not** perform any updates
 - Writers – can read and write
- Problem – allow multiple readers to read at the same time
 - Only one writer can access the shared data at the same time (i.e. no other writers or other readers are allowed)
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
 - Data set
 - Semaphore `rw_mutex` initialized to 1
 - Semaphore `mutex` initialized to 1 (to protect access to `read_count`)
 - Integer `read_count` initialized to 0

Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {  
    wait(mutex) ;  
    read_count++;  
    if (read_count == 1) /* only first reader locks rw_mutex */  
        wait(rw_mutex) ;  
    signal(mutex) ;  
  
    ...  
    /* reading dataset is performed, protected by rw_mutex */  
    /* either one writer, or multiple readers at a time  
    ...  
    wait(mutex) ;  
    read_count-- ;  
    if (read_count == 0) /* only last reader unlocks rw_mutex */  
        signal(rw_mutex) ;  
    signal(mutex) ;  
} while (true) ;
```

Readers-Writers Problem Variations

- **First** variation – no reader kept waiting unless writer has permission to use shared object
- **Second** variation – once writer is ready, it performs the write ASAP (i.e. no readers are allowed to read till after the writer gets and is done with his access)
- Both may have **starvation** leading to even more variations
- In some systems, the kernel provides a reader-writer locks

Dining-Philosophers Problem

- Philosophers spend their lives alternating between **thinking** and **eating**
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - Need both to eat, then release both when done
- In the case of 5 philosophers
 - Shared data analogy:
 - Each philosopher is a thread
 - Bowl of rice (data set)
 - Semaphore **chopstick [5]** initialized to 1



Dining-Philosophers Problem Algorithm

- The structure of Philosopher i :

```
do {  
    wait (chopstick[i] );  
    wait (chopstick[ (i + 1) % 5] );  
  
    // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```


- What is the problem with this algorithm?

Dining-Philosophers Problem Algorithm (Cont.)

- Deadlock handling
 - Allow at most 4 philosophers to be sitting simultaneously at the table of 5 chopsticks.
 - Use an asymmetric solution
 - An odd-numbered philosopher picks up first the left chopstick and then the right chopstick.
 - An even-numbered philosopher picks up first the right chopstick and then the left chopstick.
 - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section) → we may use **monitors** to implement this method.

5.8 Monitors

- The dining philosophers deadlock may be solved using monitors.
- A monitor is a high-level abstraction that provides a convenient and effective mechanism for process synchronization
- A monitor is an **abstract data type** (i.e. an **object**), internal variables only accessible by code within the procedure
- Only one process may be **active** within the monitor at a time.




5.8 Monitors

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
```

- But not powerful enough to model some synchronization schemes
- Thus we may use condition variables.




Condition Variables

- Condition variables are variables declared inside a monitor:

```
condition x, y;
```

- Two operations are allowed on a condition variable:
 - `x.wait()` – a process that invokes the operation is suspended until `x.signal()`
 - `x.signal()` – resumes one of processes (if any) that invoked `x.wait()`
- Only one thread/process can be active inside the monitor. The process that issues `x.wait()` becomes inactive, thus allowing others to be active inside the monitor.



Condition Variables – cont.

- Contrast this operation with the `signal()` operation associated with semaphores, which always affects the state of the semaphore:
 - Unlike semaphores, if there are no threads/processes that are waiting on the condition variable `x`, then calling `x.signal()` does not affect the condition variable.
 - If you recall calling `signal(my_semaphore)` increments the semaphore whether there is someone waiting on it or not.

Condition Variables Choices

- If process P invokes `x.signal()`, and process Q was suspended in `x.wait()`, what should happen next?
 - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or Q blocks waiting on another condition (i.e. immediate effect).
 - **Signal and continue** – P continues till it leaves the monitor or blocks waiting on another condition, and then Q may become active (i.e. deferred effect).
 - Both have pros and cons – language implementer can decide

Monitor solution to dining philosophers

```
monitor DiningPhilosophers
```

```
{
```

```
    enum {THINKING, HUNGRY, EATING} state[5] ;
```

```
    condition cond[5];
```

```
    void pickup (int i) {
```

```
        state[i] = HUNGRY;
```

```
        test_and_signal(i); /* if successful: my state becomes EATING + signal  
                             myself which is wasted cause I am not waiting*/
```

```
        if (state[i] != EATING) cond[i].wait(); /* test(i) did not succeed */
```

```
    }
```

```
    void putdown (int i) {
```

```
        state[i] = THINKING;
```

```
        // test left and right neighbors
```

```
        test_and_signal((i + 4) % 5);
```

```
        test_and_signal((i + 1) % 5);
```

```
    }
```

Monitor solution to Dining Philosophers (Cont.)

```
void test_and_signal(int i) {
    if ((state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING;
        cond[i].signal();
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
} /* end of monitor */
```

Monitor solution to dining philosophers (Cont.)

- Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup(i);`

`EAT`

`DiningPhilosophers.putdown(i);`

- No deadlock, but starvation is possible

Resuming Processes within a Monitor

- If several processes queued on condition `x`, and `x.signal()` executed, which should be resumed?
 - First-come-first-serve (FCFS): frequently not adequate
 - **conditional-wait** construct of the form `x.wait(c)`
 - Where `c` is **priority number**
 - Process with lowest number (highest priority) is scheduled next

Synchronization primitives used by the Linux kernel

- Prior to kernel Version 2.6, disables interrupts to implement short critical sections
- Version 2.6 and later, fully preemptive and many synchronization objects may be used within the kernel:
 - Atomic integers
 - Semaphores
 - Spinlocks
 - (with reader-writer version of semaphores and spinlocks).
- On **single-CPU** system, spinlocks replaced by **enabling and disabling kernel preemption**
 - Should be only used for short durations
- **For SMP, spinlocks are the primary tool**
 - Also should be only used for short durations.
 - A kernel thread holding a lock is not preemptable. The kernel uses a variable `preempt_count` to keep track of the number of locks a kernel thread holds.

Pthreads Synchronization

- Pthreads API is OS-independent
- It provides:
 - mutex locks
 - condition variable
- Non-portable extensions include:
 - read-write locks
 - spinlocks

Alternative approaches: OpenMP

- OpenMP is a set of compiler directives and API that support parallel programming.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

- The code contained within the **#pragma omp critical** directive is treated as a critical section and performed atomically.