

yq2021_Hw2_Problem5

October 14, 2022

Problem 5

Option 1: Finetuning a pretrained torchvision object detection model

*Some of the code comes from: https://pytorch.org/tutorials/intermediate/torchvision_tutorial.html

```
[1]: !pip install cython
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: cython in /usr/local/lib/python3.7/dist-packages (0.29.32)
```

```
[2]: !pip install pycocotools
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: pycocotools in /usr/local/lib/python3.7/dist-packages (2.0.5)
Requirement already satisfied: matplotlib>=2.1.0 in /usr/local/lib/python3.7/dist-packages (from pycocotools) (3.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from pycocotools) (1.21.6)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.1.0->pycocotools) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.1.0->pycocotools) (1.4.4)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.1.0->pycocotools) (2.8.2)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.1.0->pycocotools) (3.0.9)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from kiwisolver>=1.0.1->matplotlib>=2.1.0->pycocotools) (4.1.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->matplotlib>=2.1.0->pycocotools) (1.15.0)
```

Download the dataset and unzip it

```
[3]: !wget https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip .
```

```
--2022-10-14 20:04:37--
https://www.cis.upenn.edu/~jshi/ped_html/PennFudanPed.zip
Resolving www.cis.upenn.edu (www.cis.upenn.edu)... 158.130.69.163,
2607:f470:8:64:5ea5::d
Connecting to www.cis.upenn.edu (www.cis.upenn.edu)|158.130.69.163|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 53723336 (51M) [application/zip]
Saving to: 'PennFudanPed.zip'

PennFudanPed.zip    100%[=====>]  51.23M  10.7MB/s   in 5.9s

2022-10-14 20:04:44 (8.67 MB/s) - 'PennFudanPed.zip' saved [53723336/53723336]

--2022-10-14 20:04:44--  http://./
Resolving . (.)... failed: No address associated with hostname.
wget: unable to resolve host address '.'
FINISHED --2022-10-14 20:04:44--
Total wall clock time: 7.1s
Downloaded: 1 files, 51M in 5.9s (8.67 MB/s)
```

```
[4]: !unzip -q PennFudanPed.zip
```

Add the corresponding class for this dataset

```
[5]: import os
import numpy as np
import torch
from PIL import Image

class PennFudanDataset(torch.utils.data.Dataset):
    def __init__(self, root, transforms):
        self.root = root
        self.transforms = transforms
        # load all image files, sorting them to
        # ensure that they are aligned
        self.imgs = list(sorted(os.listdir(os.path.join(root, "PNGImages"))))
        self.masks = list(sorted(os.listdir(os.path.join(root, "PedMasks"))))

    def __getitem__(self, idx):
        # load images and masks
        img_path = os.path.join(self.root, "PNGImages", self.imgs[idx])
        mask_path = os.path.join(self.root, "PedMasks", self.masks[idx])
```

```

img = Image.open(img_path).convert("RGB")
# note that we haven't converted the mask to RGB,
# because each color corresponds to a different instance
# with 0 being background
mask = Image.open(mask_path)
# convert the PIL Image into a numpy array
mask = np.array(mask)
# instances are encoded as different colors
obj_ids = np.unique(mask)
# first id is the background, so remove it
obj_ids = obj_ids[1:]

# split the color-encoded mask into a set
# of binary masks
masks = mask == obj_ids[:, None, None]

# get bounding box coordinates for each mask
num_objs = len(obj_ids)
boxes = []
for i in range(num_objs):
    pos = np.where(masks[i])
    xmin = np.min(pos[1])
    xmax = np.max(pos[1])
    ymin = np.min(pos[0])
    ymax = np.max(pos[0])
    boxes.append([xmin, ymin, xmax, ymax])

# convert everything into a torch.Tensor
boxes = torch.as_tensor(boxes, dtype=torch.float32)
# there is only one class
labels = torch.ones((num_objs,), dtype=torch.int64)
masks = torch.as_tensor(masks, dtype=torch.uint8)

image_id = torch.tensor([idx])
area = (boxes[:, 3] - boxes[:, 1]) * (boxes[:, 2] - boxes[:, 0])
# suppose all instances are not crowd
iscrowd = torch.zeros((num_objs,), dtype=torch.int64)

target = {}
target["boxes"] = boxes
target["labels"] = labels
target["masks"] = masks
target["image_id"] = image_id
target["area"] = area
target["iscrowd"] = iscrowd

if self.transforms is not None:

```

```

        img, target = self.transforms(img, target)

    return img, target

def __len__(self):
    return len(self.imgs)

```

Check the downloaded Dataset

```
[6]: Image.open('PennFudanPed/PNGImages/FudanPed00028.png')
```

[6]:



```

[7]: test_mask = Image.open('PennFudanPed/PedMasks/FudanPed00028_mask.png')
test_mask.putpalette([
    0, 0, 0, # black background
    255, 0, 0, # index 1 is red
    0, 255, 0, # index 2 is green
    0, 0, 255, # index 3 is blue
])
test_mask

```

[7]:



Import supportive functions

```
[8]: %%shell

git clone https://github.com/pytorch/vision.git
cd vision
git checkout v0.13.0

cp references/detection/utils.py ../
cp references/detection/transforms.py ../
cp references/detection/coco_eval.py ../
cp references/detection/engine.py ../
cp references/detection/coco_utils.py ../
```

Cloning into 'vision'...

remote: Enumerating objects: 231713, done.

remote: Counting objects: 100% (5384/5384), done.

remote: Compressing objects: 100% (555/555), done.

remote: Total 231713 (delta 4960), reused 5189 (delta 4820), pack-reused

226329

Receiving objects: 100% (231713/231713), 467.54 MiB | 16.93 MiB/s, done.

Resolving deltas: 100% (210050/210050), done.

Note: checking out 'v0.13.0'.

You are in 'detached HEAD' state. You can look around, make experimental changes and commit them, and you can discard any commits you make in this state without impacting any branches by performing another checkout.

If you want to create a new branch to retain commits you create, you may do so (now or later) by using `-b` with the checkout command again. Example:

```
git checkout -b <new-branch-name>
```

HEAD is now at da3794e90 Fix all broken URLs (#6176) (#6177)

[8]:

Initialize the dataset

```
[34]: import transforms as T
      from engine import train_one_epoch, evaluate
      import utils

      def get_transform(train):
          transforms = []
          transforms.append(T.PILToTensor())
          transforms.append(T.ConvertImageDtype(torch.float))
          if train:
              transforms.append(T.RandomHorizontalFlip(0.5))
          return T.Compose(transforms)

      dataset = PennFudanDataset('PennFudanPed/', get_transform(train=True))
      dataset[0]
```

```
[34]: (tensor([[0.5804, 0.5725, 0.5608, ..., 0.8314, 0.8235, 0.8275],
               [0.5020, 0.5020, 0.4941, ..., 0.6667, 0.6784, 0.7020],
               [0.5098, 0.5137, 0.5098, ..., 0.7373, 0.7608, 0.8000],
               ...,
               [0.7255, 0.7216, 0.7176, ..., 0.8275, 0.8510, 0.8863],
               [0.7294, 0.7333, 0.7333, ..., 0.8235, 0.8588, 0.9059],
               [0.7333, 0.7451, 0.7451, ..., 0.8431, 0.8902, 0.8824]],

               [[0.3961, 0.3882, 0.3765, ..., 0.7882, 0.7804, 0.7843],
               [0.3176, 0.3176, 0.3098, ..., 0.6235, 0.6353, 0.6588],
               [0.3255, 0.3294, 0.3255, ..., 0.6941, 0.7176, 0.7569],
               ...,
               [0.7255, 0.7216, 0.7176, ..., 0.8039, 0.8275, 0.8627],
```

```

        [0.7294, 0.7333, 0.7333, ..., 0.8000, 0.8353, 0.8824],
        [0.7333, 0.7451, 0.7451, ..., 0.8196, 0.8667, 0.8588]],

[[0.3255, 0.3176, 0.3059, ..., 0.7176, 0.7098, 0.7137],
 [0.2471, 0.2471, 0.2392, ..., 0.5529, 0.5647, 0.5882],
 [0.2549, 0.2588, 0.2549, ..., 0.6235, 0.6471, 0.6863],
 ...,
 [0.7255, 0.7216, 0.7176, ..., 0.8039, 0.8275, 0.8627],
 [0.7294, 0.7333, 0.7333, ..., 0.8000, 0.8353, 0.8824],
 [0.7333, 0.7451, 0.7451, ..., 0.8196, 0.8667, 0.8588]]]),
{'boxes': tensor([[258., 181., 400., 430.],
                  [ 25., 170., 140., 485.]]),
 'labels': tensor([1, 1]),
 'masks': tensor([[[0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0],
                   ...,
                   [0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0]],
                  [[0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0],
                   ...,
                   [0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0],
                   [0, 0, 0, ..., 0, 0, 0]]], dtype=torch.uint8),
 'image_id': tensor([0]),
 'area': tensor([35358., 36225.]),
 'iscrowd': tensor([0, 0])})

```

Using Mask R-CNN to compute the instance segmentation masks

```

[35]: import torchvision
      from torchvision.models.detection.faster_rcnn import FastRCNNPredictor
      from torchvision.models.detection.mask_rcnn import MaskRCNNPredictor

      def get_model_instance_segmentation(num_classes):
          # load an instance segmentation model pre-trained on COCO
          model = torchvision.models.detection.
      ↪maskrcnn_resnet50_fpn(weights="DEFAULT")

          # get number of input features for the classifier
          in_features = model.roi_heads.box_predictor.cls_score.in_features
          # replace the pre-trained head with a new one

```

```

model.roi_heads.box_predictor = FastRCNNPredictor(in_features, num_classes)

# now get the number of input features for the mask classifier
in_features_mask = model.roi_heads.mask_predictor.conv5_mask.in_channels
hidden_layer = 256
# and replace the mask predictor with a new one
model.roi_heads.mask_predictor = MaskRCNNPredictor(in_features_mask,
                                                    hidden_layer,
                                                    num_classes)

return model

```

Testing forward() method

```

[11]: model = torchvision.models.detection.fasterrcnn_resnet50_fpn(weights="DEFAULT")
dataset = PennFudanDataset('PennFudanPed', get_transform(train=True))
data_loader = torch.utils.data.DataLoader(
    dataset, batch_size=2, shuffle=True, num_workers=4,
    collate_fn=utils.collate_fn)
# For Training
images, targets = next(iter(data_loader))
images = list(image for image in images)
targets = [{k: v for k, v in t.items()} for t in targets]
output = model(images, targets) # Returns losses and detections
# For inference
model.eval()
x = [torch.rand(3, 300, 400), torch.rand(3, 500, 400)]
predictions = model(x) # Returns predictions

```

Downloading:

"https://download.pytorch.org/models/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth"
to /root/.cache/torch/hub/checkpoints/fasterrcnn_resnet50_fpn_coco-258fb6c6.pth

0%| | 0.00/160M [00:00<?, ?B/s]

/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:566:
UserWarning: This DataLoader will create 4 worker processes in total. Our
suggested max number of worker in current system is 2, which is smaller than
what this DataLoader is going to create. Please be aware that excessive worker
creation might get DataLoader running slow or even freeze, lower the worker
number to avoid potential slowness/freeze if necessary.

cpuset_checked))

Putting Everything together

Initialize the dataset loader

```

[36]: # use our dataset and defined transformations
dataset = PennFudanDataset('PennFudanPed', get_transform(train=True))
dataset_test = PennFudanDataset('PennFudanPed', get_transform(train=False))

```



```

# split the dataset in train and test set
torch.manual_seed(3407)
indices = torch.randperm(len(dataset)).tolist()
dataset = torch.utils.data.Subset(dataset, indices[:-50])
dataset_test = torch.utils.data.Subset(dataset_test, indices[-50:])

# define training and validation data loaders
data_loader = torch.utils.data.DataLoader(
    dataset, batch_size=2, shuffle=True, num_workers=4,
    collate_fn=utils.collate_fn)

data_loader_test = torch.utils.data.DataLoader(
    dataset_test, batch_size=1, shuffle=False, num_workers=4,
    collate_fn=utils.collate_fn)

```

/usr/local/lib/python3.7/dist-packages/torch/utils/data/dataloader.py:566:
UserWarning: This DataLoader will create 4 worker processes in total. Our
suggested max number of worker in current system is 2, which is smaller than
what this DataLoader is going to create. Please be aware that excessive worker
creation might get DataLoader running slow or even freeze, lower the worker
number to avoid potential slowness/freeze if necessary.
cpuset_checked))

Initialize the parameters and optimizers

```

[37]: device = torch.device('cuda') if torch.cuda.is_available() else torch.
      ↪device('cpu')

# our dataset has two classes only - background and person
num_classes = 2

# get the model using our helper function
model = get_model_instance_segmentation(num_classes)
# move model to the right device
model.to(device)

# construct an optimizer
params = [p for p in model.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
                             momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler which decreases the learning rate by
# 10x every 3 epochs
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
                                                step_size=3,
                                                gamma=0.1)

```

Train 10 epochs using fine tuned model

```
[38]: # let's train it for 10 epochs
from torch.optim.lr_scheduler import StepLR
num_epochs = 10

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model, optimizer, data_loader, device, epoch, print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model, data_loader_test, device=device)
```

```
Epoch: [0] [ 0/60] eta: 0:01:06 lr: 0.000090 loss: 2.9246 (2.9246)
loss_classifier: 0.5471 (0.5471) loss_box_reg: 0.4399 (0.4399) loss_mask:
1.9163 (1.9163) loss_objectness: 0.0184 (0.0184) loss_rpn_box_reg: 0.0028
(0.0028) time: 1.1105 data: 0.4197 max mem: 6172
Epoch: [0] [10/60] eta: 0:00:29 lr: 0.000936 loss: 1.6280 (1.9091)
loss_classifier: 0.4320 (0.4095) loss_box_reg: 0.3044 (0.3298) loss_mask:
0.7351 (1.1391) loss_objectness: 0.0242 (0.0241) loss_rpn_box_reg: 0.0070
(0.0066) time: 0.5899 data: 0.0464 max mem: 6172
Epoch: [0] [20/60] eta: 0:00:22 lr: 0.001783 loss: 0.9292 (1.3253)
loss_classifier: 0.2026 (0.2964) loss_box_reg: 0.2471 (0.2780) loss_mask:
0.3726 (0.7260) loss_objectness: 0.0144 (0.0189) loss_rpn_box_reg: 0.0060
(0.0061) time: 0.5253 data: 0.0094 max mem: 6172
Epoch: [0] [30/60] eta: 0:00:16 lr: 0.002629 loss: 0.5777 (1.0515)
loss_classifier: 0.1136 (0.2269) loss_box_reg: 0.1805 (0.2499) loss_mask:
0.2140 (0.5528) loss_objectness: 0.0112 (0.0162) loss_rpn_box_reg: 0.0049
(0.0057) time: 0.5036 data: 0.0091 max mem: 6172
Epoch: [0] [40/60] eta: 0:00:10 lr: 0.003476 loss: 0.4099 (0.8950)
loss_classifier: 0.0613 (0.1856) loss_box_reg: 0.1561 (0.2310) loss_mask:
0.1589 (0.4589) loss_objectness: 0.0059 (0.0137) loss_rpn_box_reg: 0.0040
(0.0058) time: 0.5199 data: 0.0097 max mem: 6172
Epoch: [0] [50/60] eta: 0:00:05 lr: 0.004323 loss: 0.3021 (0.7904)
loss_classifier: 0.0498 (0.1591) loss_box_reg: 0.1269 (0.2125) loss_mask:
0.1438 (0.4007) loss_objectness: 0.0037 (0.0119) loss_rpn_box_reg: 0.0040
(0.0061) time: 0.5472 data: 0.0096 max mem: 6172
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3007 (0.7154)
loss_classifier: 0.0401 (0.1407) loss_box_reg: 0.0900 (0.1948) loss_mask:
0.1379 (0.3635) loss_objectness: 0.0012 (0.0103) loss_rpn_box_reg: 0.0038
(0.0060) time: 0.5421 data: 0.0083 max mem: 6172
Epoch: [0] Total time: 0:00:32 (0.5440 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1559 (0.1559) evaluator_time:
0.0047 (0.0047) time: 0.5019 data: 0.3401 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0908 (0.1024) evaluator_time:
```

0.0054 (0.0076) time: 0.1124 data: 0.0045 max mem: 6172
Test: Total time: 0:00:06 (0.1262 s / it)
Averaged stats: model_time: 0.0908 (0.1024) evaluator_time: 0.0054 (0.0076)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.732
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.987
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.924
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.659
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.741
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.312
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.774
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.774
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.725
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.779

IoU metric: segm

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.732
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.990
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.927
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.401
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.741
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.302
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.771
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.771
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.750
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.773

Epoch: [1] [0/60] eta: 0:01:06 lr: 0.005000 loss: 0.2624 (0.2624)
loss_classifier: 0.0314 (0.0314) loss_box_reg: 0.0726 (0.0726) loss_mask:
0.1521 (0.1521) loss_objectness: 0.0007 (0.0007) loss_rpn_box_reg: 0.0056
(0.0056) time: 1.1027 data: 0.4638 max mem: 6172
Epoch: [1] [10/60] eta: 0:00:30 lr: 0.005000 loss: 0.2346 (0.2546)
loss_classifier: 0.0314 (0.0329) loss_box_reg: 0.0722 (0.0830) loss_mask:
0.1407 (0.1327) loss_objectness: 0.0007 (0.0011) loss_rpn_box_reg: 0.0035
(0.0049) time: 0.6024 data: 0.0496 max mem: 6172
Epoch: [1] [20/60] eta: 0:00:22 lr: 0.005000 loss: 0.2323 (0.2472)
loss_classifier: 0.0291 (0.0319) loss_box_reg: 0.0686 (0.0736) loss_mask:
0.1326 (0.1352) loss_objectness: 0.0009 (0.0017) loss_rpn_box_reg: 0.0038
(0.0050) time: 0.5373 data: 0.0083 max mem: 6172
Epoch: [1] [30/60] eta: 0:00:16 lr: 0.005000 loss: 0.2316 (0.2590)
loss_classifier: 0.0291 (0.0351) loss_box_reg: 0.0612 (0.0735) loss_mask:
0.1423 (0.1437) loss_objectness: 0.0012 (0.0017) loss_rpn_box_reg: 0.0039

```

(0.0050) time: 0.5443 data: 0.0085 max mem: 6172
Epoch: [1] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.2955 (0.2642)
loss_classifier: 0.0349 (0.0364) loss_box_reg: 0.0645 (0.0747) loss_mask:
0.1541 (0.1464) loss_objectness: 0.0011 (0.0018) loss_rpn_box_reg: 0.0034
(0.0049) time: 0.5501 data: 0.0089 max mem: 6172
Epoch: [1] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2238 (0.2547)
loss_classifier: 0.0312 (0.0346) loss_box_reg: 0.0597 (0.0714) loss_mask:
0.1249 (0.1423) loss_objectness: 0.0008 (0.0016) loss_rpn_box_reg: 0.0034
(0.0048) time: 0.5248 data: 0.0093 max mem: 6172
Epoch: [1] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2432 (0.2613)
loss_classifier: 0.0368 (0.0360) loss_box_reg: 0.0634 (0.0745) loss_mask:
0.1249 (0.1441) loss_objectness: 0.0004 (0.0016) loss_rpn_box_reg: 0.0038
(0.0051) time: 0.5392 data: 0.0093 max mem: 6172
Epoch: [1] Total time: 0:00:33 (0.5559 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.1455 (0.1455) evaluator_time:
0.0049 (0.0049) time: 0.4891 data: 0.3375 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0896 (0.1025) evaluator_time:
0.0048 (0.0064) time: 0.1104 data: 0.0044 max mem: 6172
Test: Total time: 0:00:06 (0.1251 s / it)
Averaged stats: model_time: 0.0896 (0.1025) evaluator_time: 0.0048 (0.0064)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.759
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.995
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.926
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.715
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.765
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.319
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.798
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.798
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.750
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.803
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.752
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.995
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.919
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.601
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.757
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.315
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.785

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.785
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.775
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.786
Epoch: [2] [ 0/60] eta: 0:01:05 lr: 0.005000 loss: 0.2863 (0.2863)
loss_classifier: 0.0435 (0.0435) loss_box_reg: 0.0752 (0.0752) loss_mask:
0.1568 (0.1568) loss_objectness: 0.0024 (0.0024) loss_rpn_box_reg: 0.0084
(0.0084) time: 1.0846 data: 0.4351 max mem: 6172
Epoch: [2] [10/60] eta: 0:00:28 lr: 0.005000 loss: 0.1999 (0.2222)
loss_classifier: 0.0240 (0.0306) loss_box_reg: 0.0514 (0.0565) loss_mask:
0.1246 (0.1297) loss_objectness: 0.0007 (0.0016) loss_rpn_box_reg: 0.0037
(0.0039) time: 0.5775 data: 0.0464 max mem: 6172
Epoch: [2] [20/60] eta: 0:00:22 lr: 0.005000 loss: 0.2067 (0.2410)
loss_classifier: 0.0290 (0.0347) loss_box_reg: 0.0558 (0.0692) loss_mask:
0.1211 (0.1313) loss_objectness: 0.0007 (0.0014) loss_rpn_box_reg: 0.0031
(0.0044) time: 0.5482 data: 0.0088 max mem: 6172
Epoch: [2] [30/60] eta: 0:00:17 lr: 0.005000 loss: 0.2004 (0.2266)
loss_classifier: 0.0290 (0.0320) loss_box_reg: 0.0496 (0.0622) loss_mask:
0.1118 (0.1271) loss_objectness: 0.0006 (0.0013) loss_rpn_box_reg: 0.0033
(0.0041) time: 0.5614 data: 0.0092 max mem: 6172
Epoch: [2] [40/60] eta: 0:00:11 lr: 0.005000 loss: 0.1869 (0.2231)
loss_classifier: 0.0244 (0.0310) loss_box_reg: 0.0474 (0.0609) loss_mask:
0.1169 (0.1262) loss_objectness: 0.0006 (0.0013) loss_rpn_box_reg: 0.0028
(0.0038) time: 0.5329 data: 0.0088 max mem: 6172
Epoch: [2] [50/60] eta: 0:00:05 lr: 0.005000 loss: 0.2214 (0.2263)
loss_classifier: 0.0291 (0.0316) loss_box_reg: 0.0523 (0.0619) loss_mask:
0.1269 (0.1277) loss_objectness: 0.0003 (0.0014) loss_rpn_box_reg: 0.0025
(0.0037) time: 0.5538 data: 0.0095 max mem: 6172
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.2058 (0.2215)
loss_classifier: 0.0298 (0.0302) loss_box_reg: 0.0473 (0.0595) loss_mask:
0.1129 (0.1266) loss_objectness: 0.0004 (0.0013) loss_rpn_box_reg: 0.0029
(0.0040) time: 0.5570 data: 0.0090 max mem: 6172
Epoch: [2] Total time: 0:00:33 (0.5589 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.1596 (0.1596) evaluator_time:
0.0048 (0.0048) time: 0.4975 data: 0.3318 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0908 (0.1024) evaluator_time:
0.0048 (0.0056) time: 0.1097 data: 0.0044 max mem: 6172
Test: Total time: 0:00:06 (0.1243 s / it)
Averaged stats: model_time: 0.0908 (0.1024) evaluator_time: 0.0048 (0.0056)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.809
Average Precision (AP) @[ IoU=0.50 | area=  all | maxDets=100 ] = 0.991

```

Average Precision	(AP) @[IoU=0.75	area= all maxDets=100]	= 0.946
Average Precision	(AP) @[IoU=0.50:0.95	area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95	area=medium maxDets=100]	= 0.693
Average Precision	(AP) @[IoU=0.50:0.95	area= large maxDets=100]	= 0.821
Average Recall	(AR) @[IoU=0.50:0.95	area= all maxDets= 1]	= 0.345
Average Recall	(AR) @[IoU=0.50:0.95	area= all maxDets= 10]	= 0.845
Average Recall	(AR) @[IoU=0.50:0.95	area= all maxDets=100]	= 0.845
Average Recall	(AR) @[IoU=0.50:0.95	area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95	area=medium maxDets=100]	= 0.750
Average Recall	(AR) @[IoU=0.50:0.95	area= large maxDets=100]	= 0.854
IoU metric: segm			
Average Precision	(AP) @[IoU=0.50:0.95	area= all maxDets=100]	= 0.774
Average Precision	(AP) @[IoU=0.50	area= all maxDets=100]	= 0.983
Average Precision	(AP) @[IoU=0.75	area= all maxDets=100]	= 0.939
Average Precision	(AP) @[IoU=0.50:0.95	area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95	area=medium maxDets=100]	= 0.571
Average Precision	(AP) @[IoU=0.50:0.95	area= large maxDets=100]	= 0.782
Average Recall	(AR) @[IoU=0.50:0.95	area= all maxDets= 1]	= 0.325
Average Recall	(AR) @[IoU=0.50:0.95	area= all maxDets= 10]	= 0.802
Average Recall	(AR) @[IoU=0.50:0.95	area= all maxDets=100]	= 0.802
Average Recall	(AR) @[IoU=0.50:0.95	area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95	area=medium maxDets=100]	= 0.767
Average Recall	(AR) @[IoU=0.50:0.95	area= large maxDets=100]	= 0.805
Epoch: [3] [0/60] eta: 0:00:52 lr: 0.000500 loss: 0.1661 (0.1661)			
loss_classifier: 0.0232 (0.0232) loss_box_reg: 0.0222 (0.0222) loss_mask:			
0.1182 (0.1182) loss_objectness: 0.0016 (0.0016) loss_rpn_box_reg: 0.0009			
(0.0009) time: 0.8757 data: 0.3491 max mem: 6172			
Epoch: [3] [10/60] eta: 0:00:29 lr: 0.000500 loss: 0.1787 (0.1937)			
loss_classifier: 0.0253 (0.0262) loss_box_reg: 0.0380 (0.0444) loss_mask:			
0.1163 (0.1194) loss_objectness: 0.0003 (0.0006) loss_rpn_box_reg: 0.0020			
(0.0031) time: 0.5821 data: 0.0396 max mem: 6172			
Epoch: [3] [20/60] eta: 0:00:22 lr: 0.000500 loss: 0.1787 (0.1896)			
loss_classifier: 0.0243 (0.0246) loss_box_reg: 0.0370 (0.0405) loss_mask:			
0.1129 (0.1205) loss_objectness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0018			
(0.0031) time: 0.5398 data: 0.0089 max mem: 6172			
Epoch: [3] [30/60] eta: 0:00:16 lr: 0.000500 loss: 0.1646 (0.1820)			
loss_classifier: 0.0214 (0.0242) loss_box_reg: 0.0307 (0.0376) loss_mask:			
0.1115 (0.1162) loss_objectness: 0.0004 (0.0010) loss_rpn_box_reg: 0.0017			
(0.0029) time: 0.5303 data: 0.0088 max mem: 6172			
Epoch: [3] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1763 (0.1848)			
loss_classifier: 0.0228 (0.0251) loss_box_reg: 0.0346 (0.0388) loss_mask:			
0.1107 (0.1166) loss_objectness: 0.0004 (0.0010) loss_rpn_box_reg: 0.0023			
(0.0032) time: 0.5671 data: 0.0089 max mem: 6172			
Epoch: [3] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1824 (0.1883)			
loss_classifier: 0.0236 (0.0264) loss_box_reg: 0.0352 (0.0402) loss_mask:			
0.1127 (0.1177) loss_objectness: 0.0002 (0.0009) loss_rpn_box_reg: 0.0023			
(0.0032) time: 0.5709 data: 0.0091 max mem: 6172			
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1839 (0.1872)			

```

loss_classifier: 0.0247 (0.0261) loss_box_reg: 0.0382 (0.0399) loss_mask:
0.1214 (0.1174) loss_objectness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0016
(0.0030) time: 0.5565 data: 0.0085 max mem: 6172
Epoch: [3] Total time: 0:00:33 (0.5621 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1518 (0.1518) evaluator_time:
0.0048 (0.0048) time: 0.5074 data: 0.3496 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0896 (0.1026) evaluator_time:
0.0050 (0.0055) time: 0.1097 data: 0.0043 max mem: 6172
Test: Total time: 0:00:06 (0.1242 s / it)
Averaged stats: model_time: 0.0896 (0.1026) evaluator_time: 0.0050 (0.0055)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.826
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.990
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.943
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.716
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.838
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.351
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.862
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.862
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.869
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.772
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.983
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.939
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.529
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.781
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.322
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.803
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.775
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.806
Epoch: [4] [ 0/60] eta: 0:00:57 lr: 0.000500 loss: 0.1606 (0.1606)
loss_classifier: 0.0258 (0.0258) loss_box_reg: 0.0292 (0.0292) loss_mask:
0.1013 (0.1013) loss_objectness: 0.0017 (0.0017) loss_rpn_box_reg: 0.0025
(0.0025) time: 0.9666 data: 0.3818 max mem: 6172
Epoch: [4] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.1779 (0.1839)
loss_classifier: 0.0258 (0.0271) loss_box_reg: 0.0298 (0.0369) loss_mask:

```

```

0.1033 (0.1154) loss_objectness: 0.0006 (0.0012) loss_rpn_box_reg: 0.0025
(0.0032) time: 0.6027 data: 0.0425 max mem: 6172
Epoch: [4] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1670 (0.1870)
loss_classifier: 0.0255 (0.0276) loss_box_reg: 0.0317 (0.0401) loss_mask:
0.1057 (0.1152) loss_objectness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0025
(0.0033) time: 0.5679 data: 0.0088 max mem: 6172
Epoch: [4] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1663 (0.1857)
loss_classifier: 0.0238 (0.0272) loss_box_reg: 0.0345 (0.0398) loss_mask:
0.1057 (0.1147) loss_objectness: 0.0002 (0.0011) loss_rpn_box_reg: 0.0021
(0.0029) time: 0.5550 data: 0.0101 max mem: 6172
Epoch: [4] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1715 (0.1837)
loss_classifier: 0.0224 (0.0260) loss_box_reg: 0.0359 (0.0387) loss_mask:
0.1079 (0.1150) loss_objectness: 0.0003 (0.0010) loss_rpn_box_reg: 0.0016
(0.0031) time: 0.5502 data: 0.0098 max mem: 6172
Epoch: [4] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1681 (0.1843)
loss_classifier: 0.0224 (0.0260) loss_box_reg: 0.0309 (0.0391) loss_mask:
0.1064 (0.1153) loss_objectness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0019
(0.0029) time: 0.5622 data: 0.0088 max mem: 6172
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1510 (0.1812)
loss_classifier: 0.0204 (0.0258) loss_box_reg: 0.0306 (0.0382) loss_mask:
0.1002 (0.1136) loss_objectness: 0.0002 (0.0009) loss_rpn_box_reg: 0.0017
(0.0028) time: 0.5623 data: 0.0087 max mem: 6172
Epoch: [4] Total time: 0:00:34 (0.5683 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1464 (0.1464) evaluator_time:
0.0050 (0.0050) time: 0.5019 data: 0.3493 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0900 (0.1035) evaluator_time:
0.0045 (0.0056) time: 0.1110 data: 0.0045 max mem: 6172
Test: Total time: 0:00:06 (0.1254 s / it)
Averaged stats: model_time: 0.0900 (0.1035) evaluator_time: 0.0045 (0.0056)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.828
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.990
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.942
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.729
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.841
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.350
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.865
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.865
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.792
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.872

```



```

IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.775
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.982
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.939
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.524
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.782
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.324
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.805
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.805
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.783
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.808
Epoch: [5] [ 0/60] eta: 0:01:01 lr: 0.000500 loss: 0.2103 (0.2103)
loss_classifier: 0.0309 (0.0309) loss_box_reg: 0.0322 (0.0322) loss_mask:
0.1454 (0.1454) loss_objectness: 0.0006 (0.0006) loss_rpn_box_reg: 0.0011
(0.0011) time: 1.0329 data: 0.4769 max mem: 6172
Epoch: [5] [10/60] eta: 0:00:30 lr: 0.000500 loss: 0.1709 (0.1848)
loss_classifier: 0.0228 (0.0266) loss_box_reg: 0.0322 (0.0386) loss_mask:
0.1075 (0.1164) loss_objectness: 0.0002 (0.0003) loss_rpn_box_reg: 0.0031
(0.0028) time: 0.6062 data: 0.0528 max mem: 6172
Epoch: [5] [20/60] eta: 0:00:23 lr: 0.000500 loss: 0.1709 (0.1827)
loss_classifier: 0.0227 (0.0263) loss_box_reg: 0.0290 (0.0365) loss_mask:
0.1075 (0.1166) loss_objectness: 0.0002 (0.0006) loss_rpn_box_reg: 0.0021
(0.0027) time: 0.5621 data: 0.0101 max mem: 6172
Epoch: [5] [30/60] eta: 0:00:17 lr: 0.000500 loss: 0.1601 (0.1798)
loss_classifier: 0.0163 (0.0260) loss_box_reg: 0.0286 (0.0356) loss_mask:
0.1075 (0.1149) loss_objectness: 0.0003 (0.0007) loss_rpn_box_reg: 0.0018
(0.0026) time: 0.5591 data: 0.0094 max mem: 6172
Epoch: [5] [40/60] eta: 0:00:11 lr: 0.000500 loss: 0.1510 (0.1769)
loss_classifier: 0.0177 (0.0253) loss_box_reg: 0.0233 (0.0345) loss_mask:
0.1043 (0.1136) loss_objectness: 0.0002 (0.0008) loss_rpn_box_reg: 0.0018
(0.0027) time: 0.5461 data: 0.0090 max mem: 6172
Epoch: [5] [50/60] eta: 0:00:05 lr: 0.000500 loss: 0.1621 (0.1816)
loss_classifier: 0.0268 (0.0264) loss_box_reg: 0.0290 (0.0368) loss_mask:
0.1043 (0.1145) loss_objectness: 0.0002 (0.0008) loss_rpn_box_reg: 0.0023
(0.0030) time: 0.5581 data: 0.0091 max mem: 6172
Epoch: [5] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.1751 (0.1785)
loss_classifier: 0.0267 (0.0255) loss_box_reg: 0.0371 (0.0359) loss_mask:
0.1065 (0.1136) loss_objectness: 0.0002 (0.0007) loss_rpn_box_reg: 0.0023
(0.0029) time: 0.5816 data: 0.0088 max mem: 6172
Epoch: [5] Total time: 0:00:34 (0.5735 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.1676 (0.1676) evaluator_time:
0.0041 (0.0041) time: 0.4806 data: 0.3075 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0908 (0.1044) evaluator_time:
0.0039 (0.0054) time: 0.1117 data: 0.0048 max mem: 6172

```

Test: Total time: 0:00:06 (0.1256 s / it)
Averaged stats: model_time: 0.0908 (0.1044) evaluator_time: 0.0039 (0.0054)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.831
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.990
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.938
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.738
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.841
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.350
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.863
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.863
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.800
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.869

IoU metric: segm

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.774
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.982
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.919
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.513
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.783
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.324
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.803
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.803
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.775
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.806

Epoch: [6] [0/60] eta: 0:00:56 lr: 0.000050 loss: 0.1186 (0.1186)
loss_classifier: 0.0127 (0.0127) loss_box_reg: 0.0150 (0.0150) loss_mask:
0.0894 (0.0894) loss_objectness: 0.0002 (0.0002) loss_rpn_box_reg: 0.0013
(0.0013) time: 0.9372 data: 0.3950 max mem: 6172
Epoch: [6] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1873 (0.1908)
loss_classifier: 0.0281 (0.0287) loss_box_reg: 0.0329 (0.0383) loss_mask:
0.1094 (0.1185) loss_objectness: 0.0010 (0.0021) loss_rpn_box_reg: 0.0030
(0.0032) time: 0.5878 data: 0.0437 max mem: 6172
Epoch: [6] [20/60] eta: 0:00:22 lr: 0.000050 loss: 0.1779 (0.1852)
loss_classifier: 0.0240 (0.0269) loss_box_reg: 0.0329 (0.0353) loss_mask:
0.1094 (0.1183) loss_objectness: 0.0003 (0.0018) loss_rpn_box_reg: 0.0023
(0.0028) time: 0.5484 data: 0.0089 max mem: 6172
Epoch: [6] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1698 (0.1877)
loss_classifier: 0.0222 (0.0269) loss_box_reg: 0.0353 (0.0379) loss_mask:
0.1087 (0.1186) loss_objectness: 0.0002 (0.0014) loss_rpn_box_reg: 0.0018
(0.0029) time: 0.5784 data: 0.0091 max mem: 6172

```

Epoch: [6] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1514 (0.1779)
loss_classifier: 0.0188 (0.0248) loss_box_reg: 0.0255 (0.0339) loss_mask:
0.1048 (0.1155) loss_objectness: 0.0003 (0.0012) loss_rpn_box_reg: 0.0016
(0.0026) time: 0.5715 data: 0.0086 max mem: 6172
Epoch: [6] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1514 (0.1776)
loss_classifier: 0.0188 (0.0246) loss_box_reg: 0.0261 (0.0345) loss_mask:
0.1033 (0.1145) loss_objectness: 0.0003 (0.0012) loss_rpn_box_reg: 0.0021
(0.0027) time: 0.5737 data: 0.0084 max mem: 6172
Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1619 (0.1754)
loss_classifier: 0.0217 (0.0243) loss_box_reg: 0.0296 (0.0341) loss_mask:
0.1030 (0.1133) loss_objectness: 0.0002 (0.0010) loss_rpn_box_reg: 0.0026
(0.0027) time: 0.5817 data: 0.0083 max mem: 6172
Epoch: [6] Total time: 0:00:34 (0.5782 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:24 model_time: 0.1573 (0.1573) evaluator_time:
0.0041 (0.0041) time: 0.4806 data: 0.3180 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0929 (0.1038) evaluator_time:
0.0041 (0.0055) time: 0.1111 data: 0.0041 max mem: 6172
Test: Total time: 0:00:06 (0.1253 s / it)
Averaged stats: model_time: 0.0929 (0.1038) evaluator_time: 0.0041 (0.0055)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.833
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.931
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.735
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.843
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.352
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.865
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.865
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.800
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.872
IoU metric: segm
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.779
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.982
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.920
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.518
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.789
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.327
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.807
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.807

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.792
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.808
Epoch: [7] [ 0/60] eta: 0:00:53 lr: 0.000050 loss: 0.1325 (0.1325)
loss_classifier: 0.0098 (0.0098) loss_box_reg: 0.0160 (0.0160) loss_mask:
0.1062 (0.1062) loss_objectness: 0.0001 (0.0001) loss_rpn_box_reg: 0.0005
(0.0005) time: 0.8950 data: 0.3721 max mem: 6172
Epoch: [7] [10/60] eta: 0:00:29 lr: 0.000050 loss: 0.1486 (0.1647)
loss_classifier: 0.0231 (0.0219) loss_box_reg: 0.0298 (0.0309) loss_mask:
0.1062 (0.1096) loss_objectness: 0.0002 (0.0004) loss_rpn_box_reg: 0.0015
(0.0019) time: 0.5890 data: 0.0418 max mem: 6172
Epoch: [7] [20/60] eta: 0:00:23 lr: 0.000050 loss: 0.1624 (0.1701)
loss_classifier: 0.0243 (0.0229) loss_box_reg: 0.0340 (0.0325) loss_mask:
0.1082 (0.1116) loss_objectness: 0.0002 (0.0004) loss_rpn_box_reg: 0.0023
(0.0026) time: 0.5605 data: 0.0090 max mem: 6172
Epoch: [7] [30/60] eta: 0:00:17 lr: 0.000050 loss: 0.1565 (0.1678)
loss_classifier: 0.0207 (0.0237) loss_box_reg: 0.0290 (0.0318) loss_mask:
0.1078 (0.1092) loss_objectness: 0.0002 (0.0006) loss_rpn_box_reg: 0.0023
(0.0025) time: 0.5627 data: 0.0089 max mem: 6172
Epoch: [7] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1448 (0.1695)
loss_classifier: 0.0163 (0.0235) loss_box_reg: 0.0218 (0.0328) loss_mask:
0.1101 (0.1102) loss_objectness: 0.0003 (0.0005) loss_rpn_box_reg: 0.0019
(0.0026) time: 0.5715 data: 0.0087 max mem: 6172
Epoch: [7] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1841 (0.1749)
loss_classifier: 0.0212 (0.0241) loss_box_reg: 0.0340 (0.0343) loss_mask:
0.1113 (0.1129) loss_objectness: 0.0002 (0.0008) loss_rpn_box_reg: 0.0027
(0.0027) time: 0.5844 data: 0.0089 max mem: 6172
Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1798 (0.1758)
loss_classifier: 0.0232 (0.0243) loss_box_reg: 0.0243 (0.0343) loss_mask:
0.1113 (0.1136) loss_objectness: 0.0003 (0.0008) loss_rpn_box_reg: 0.0024
(0.0027) time: 0.5675 data: 0.0085 max mem: 6172
Epoch: [7] Total time: 0:00:34 (0.5769 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:23 model_time: 0.1528 (0.1528) evaluator_time:
0.0040 (0.0040) time: 0.4673 data: 0.3091 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0905 (0.1036) evaluator_time:
0.0039 (0.0054) time: 0.1117 data: 0.0048 max mem: 6172
Test: Total time: 0:00:06 (0.1249 s / it)
Averaged stats: model_time: 0.0905 (0.1036) evaluator_time: 0.0039 (0.0054)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.831
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.930

```

Average Precision	(AP)	@[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP)	@[IoU=0.50:0.95 area=medium maxDets=100]	= 0.737
Average Precision	(AP)	@[IoU=0.50:0.95 area= large maxDets=100]	= 0.842
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets= 1]	= 0.352
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets= 10]	= 0.866
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets=100]	= 0.866
Average Recall	(AR)	@[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR)	@[IoU=0.50:0.95 area=medium maxDets=100]	= 0.808
Average Recall	(AR)	@[IoU=0.50:0.95 area= large maxDets=100]	= 0.872

IoU metric: segm

Average Precision	(AP)	@[IoU=0.50:0.95 area= all maxDets=100]	= 0.774
Average Precision	(AP)	@[IoU=0.50 area= all maxDets=100]	= 0.982
Average Precision	(AP)	@[IoU=0.75 area= all maxDets=100]	= 0.929
Average Precision	(AP)	@[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP)	@[IoU=0.50:0.95 area=medium maxDets=100]	= 0.523
Average Precision	(AP)	@[IoU=0.50:0.95 area= large maxDets=100]	= 0.782
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets= 1]	= 0.325
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets= 10]	= 0.805
Average Recall	(AR)	@[IoU=0.50:0.95 area= all maxDets=100]	= 0.805
Average Recall	(AR)	@[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR)	@[IoU=0.50:0.95 area=medium maxDets=100]	= 0.792
Average Recall	(AR)	@[IoU=0.50:0.95 area= large maxDets=100]	= 0.806

Epoch: [8] [0/60] eta: 0:00:57 lr: 0.000050 loss: 0.1174 (0.1174)
loss_classifier: 0.0051 (0.0051) loss_box_reg: 0.0128 (0.0128) loss_mask:
0.0958 (0.0958) loss_objectness: 0.0002 (0.0002) loss_rpn_box_reg: 0.0035
(0.0035) time: 0.9555 data: 0.4094 max mem: 6172

Epoch: [8] [10/60] eta: 0:00:28 lr: 0.000050 loss: 0.1461 (0.1516)
loss_classifier: 0.0171 (0.0188) loss_box_reg: 0.0255 (0.0247) loss_mask:
0.0995 (0.1058) loss_objectness: 0.0002 (0.0003) loss_rpn_box_reg: 0.0020
(0.0020) time: 0.5746 data: 0.0447 max mem: 6172

Epoch: [8] [20/60] eta: 0:00:23 lr: 0.000050 loss: 0.1564 (0.1661)
loss_classifier: 0.0218 (0.0225) loss_box_reg: 0.0257 (0.0306) loss_mask:
0.1035 (0.1096) loss_objectness: 0.0003 (0.0011) loss_rpn_box_reg: 0.0014
(0.0023) time: 0.5605 data: 0.0085 max mem: 6172

Epoch: [8] [30/60] eta: 0:00:16 lr: 0.000050 loss: 0.1774 (0.1734)
loss_classifier: 0.0242 (0.0239) loss_box_reg: 0.0310 (0.0341) loss_mask:
0.1058 (0.1120) loss_objectness: 0.0005 (0.0010) loss_rpn_box_reg: 0.0016
(0.0024) time: 0.5613 data: 0.0106 max mem: 6172

Epoch: [8] [40/60] eta: 0:00:11 lr: 0.000050 loss: 0.1696 (0.1753)
loss_classifier: 0.0210 (0.0246) loss_box_reg: 0.0310 (0.0353) loss_mask:
0.0984 (0.1120) loss_objectness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0022
(0.0026) time: 0.5528 data: 0.0108 max mem: 6172

Epoch: [8] [50/60] eta: 0:00:05 lr: 0.000050 loss: 0.1619 (0.1748)
loss_classifier: 0.0223 (0.0244) loss_box_reg: 0.0259 (0.0344) loss_mask:
0.1105 (0.1125) loss_objectness: 0.0003 (0.0010) loss_rpn_box_reg: 0.0019
(0.0025) time: 0.5794 data: 0.0095 max mem: 6172

Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.1636 (0.1756)
loss_classifier: 0.0237 (0.0246) loss_box_reg: 0.0290 (0.0349) loss_mask:

```

0.1110 (0.1126)  loss_objectness: 0.0004 (0.0010)  loss_rpn_box_reg: 0.0019
(0.0025)  time: 0.5968  data: 0.0092  max mem: 6172
Epoch: [8] Total time: 0:00:34 (0.5783 s / it)
creating index...
index created!
Test: [ 0/50]  eta: 0:00:22  model_time: 0.1608 (0.1608)  evaluator_time:
0.0040 (0.0040)  time: 0.4594  data: 0.2916  max mem: 6172
Test: [49/50]  eta: 0:00:00  model_time: 0.0906 (0.1043)  evaluator_time:
0.0041 (0.0054)  time: 0.1112  data: 0.0044  max mem: 6172
Test: Total time: 0:00:06 (0.1249 s / it)
Averaged stats: model_time: 0.0906 (0.1043)  evaluator_time: 0.0041 (0.0054)
Accumulating evaluation results...
DONE (t=0.01s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.830
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.991
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.930
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.732
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.841
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.352
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.864
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.864
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.800
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.870
IoU metric: segm
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.774
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.982
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.929
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.521
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.782
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.324
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.803
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.803
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.783
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.805
Epoch: [9] [ 0/60]  eta: 0:01:10  lr: 0.000005  loss: 0.1382 (0.1382)
loss_classifier: 0.0135 (0.0135)  loss_box_reg: 0.0210 (0.0210)  loss_mask:
0.0982 (0.0982)  loss_objectness: 0.0040 (0.0040)  loss_rpn_box_reg: 0.0015
(0.0015)  time: 1.1709  data: 0.4315  max mem: 6172
Epoch: [9] [10/60]  eta: 0:00:30  lr: 0.000005  loss: 0.1464 (0.1678)
loss_classifier: 0.0203 (0.0226)  loss_box_reg: 0.0240 (0.0317)  loss_mask:
0.1024 (0.1102)  loss_objectness: 0.0005 (0.0008)  loss_rpn_box_reg: 0.0016

```

```

(0.0026) time: 0.6173 data: 0.0455 max mem: 6172
Epoch: [9] [20/60] eta: 0:00:23 lr: 0.000005 loss: 0.1594 (0.1706)
loss_classifier: 0.0211 (0.0238) loss_box_reg: 0.0240 (0.0302) loss_mask:
0.1046 (0.1135) loss_objectness: 0.0001 (0.0008) loss_rpn_box_reg: 0.0016
(0.0023) time: 0.5627 data: 0.0079 max mem: 6172
Epoch: [9] [30/60] eta: 0:00:17 lr: 0.000005 loss: 0.1813 (0.1764)
loss_classifier: 0.0247 (0.0248) loss_box_reg: 0.0349 (0.0341) loss_mask:
0.1204 (0.1142) loss_objectness: 0.0001 (0.0009) loss_rpn_box_reg: 0.0016
(0.0024) time: 0.5751 data: 0.0088 max mem: 6172
Epoch: [9] [40/60] eta: 0:00:11 lr: 0.000005 loss: 0.1865 (0.1767)
loss_classifier: 0.0243 (0.0247) loss_box_reg: 0.0401 (0.0346) loss_mask:
0.1171 (0.1141) loss_objectness: 0.0003 (0.0008) loss_rpn_box_reg: 0.0018
(0.0025) time: 0.5759 data: 0.0089 max mem: 6172
Epoch: [9] [50/60] eta: 0:00:05 lr: 0.000005 loss: 0.1659 (0.1773)
loss_classifier: 0.0234 (0.0250) loss_box_reg: 0.0276 (0.0347) loss_mask:
0.1147 (0.1142) loss_objectness: 0.0003 (0.0009) loss_rpn_box_reg: 0.0023
(0.0025) time: 0.5656 data: 0.0091 max mem: 6172
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.1567 (0.1738)
loss_classifier: 0.0176 (0.0242) loss_box_reg: 0.0232 (0.0337) loss_mask:
0.1008 (0.1127) loss_objectness: 0.0003 (0.0008) loss_rpn_box_reg: 0.0020
(0.0025) time: 0.5521 data: 0.0086 max mem: 6172
Epoch: [9] Total time: 0:00:34 (0.5781 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:25 model_time: 0.1645 (0.1645) evaluator_time:
0.0046 (0.0046) time: 0.5018 data: 0.3316 max mem: 6172
Test: [49/50] eta: 0:00:00 model_time: 0.0906 (0.1047) evaluator_time:
0.0041 (0.0055) time: 0.1121 data: 0.0045 max mem: 6172
Test: Total time: 0:00:06 (0.1264 s / it)
Averaged stats: model_time: 0.0906 (0.1047) evaluator_time: 0.0041 (0.0055)
Accumulating evaluation results...
DONE (t=0.02s).
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.830
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.991
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.930
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.732
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.841
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.352
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.865
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.865
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.800
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.871
IoU metric: segm

```

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.772
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.982
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.929
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.521
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.779
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.323
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.801
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.801
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.783
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.803

Check the output of the trained model

```
[41]: # pick random one image from the test set
import random
index = random.randint(0, 50)
img, _ = dataset_test[index]
# put the model in evaluation mode
model.eval()
with torch.no_grad():
    prediction = model([img.to(device)])
```

```
[42]: # check the predicted result
prediction
```

```
[42]: [{'boxes': tensor([[257.1166, 64.9268, 413.6607, 351.0188],
                      [189.0505, 96.6888, 210.3089, 171.2245],
                      [303.5939, 60.2320, 388.4765, 314.0001]]), device='cuda:0'),
      'labels': tensor([1, 1, 1], device='cuda:0'),
      'scores': tensor([0.9975, 0.9773, 0.2118], device='cuda:0'),
      'masks': tensor([[[[0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.],
                        ...,
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.]]],
                      [[0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.],
                        ...,
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.],
                        [0., 0., 0., ..., 0., 0., 0.]]])]
```



```
[[[0., 0., 0., ..., 0., 0., 0.],
  [0., 0., 0., ..., 0., 0., 0.],
  [0., 0., 0., ..., 0., 0., 0.],
  ...,
  [0., 0., 0., ..., 0., 0., 0.],
  [0., 0., 0., ..., 0., 0., 0.],
  [0., 0., 0., ..., 0., 0., 0.] ]], device='cuda:0')}]
```

Reformat the image and check the result

```
[43]: Image.fromarray(img.mul(255).permute(1, 2, 0).byte().numpy())
```

[43]:



Check the segmentation mask


```
[44]: Image.fromarray(prediction[0]['masks'][0, 0].mul(255).byte().cpu().numpy())
```

[44]:



```
[45]: Image.fromarray(prediction[0]['masks'][1, 0].mul(255).byte().cpu().numpy())
```

```
[45]:
```



Option 2: Add another backbone

```
[46]: import torchvision
      from torchvision.models.detection import FasterRCNN
      from torchvision.models.detection.rpn import AnchorGenerator

      # load a pre-trained model for classification and return
      # only the features
      backbone = torchvision.models.mobilenet_v2(weights="DEFAULT").features
      # FasterRCNN needs to know the number of
      # output channels in a backbone. For mobilenet_v2, it's 1280
      # so we need to add it here
      backbone.out_channels = 1280

      # let's make the RPN generate 5 x 3 anchors per spatial
      # location, with 5 different sizes and 3 different aspect
      # ratios. We have a Tuple[Tuple[int]] because each feature
      # map could potentially have different sizes and
```

```

# aspect ratios
anchor_generator = AnchorGenerator(sizes=((32, 64, 128, 256, 512),),
                                   aspect_ratios=((0.5, 1.0, 2.0),))

# let's define what are the feature maps that we will
# use to perform the region of interest cropping, as well as
# the size of the crop after rescaling.
# if your backbone returns a Tensor, featmap_names is expected to
# be [0]. More generally, the backbone should return an
# OrderedDict[Tensor], and in featmap_names you can choose which
# feature maps to use.
roi_pooler = torchvision.ops.MultiScaleRoIAlign(featmap_names=['0'],
                                                output_size=7,
                                                sampling_ratio=2)

# put the pieces together inside a FasterRCNN model
model_option_2 = FasterRCNN(backbone,
                             num_classes=2,
                             rpn_anchor_generator=anchor_generator,
                             box_roi_pool=roi_pooler)

device = torch.device('cuda') if torch.cuda.is_available() else torch.
    ↪device('cpu')

# move model to the right device
# Changes
model_option_2.to(device)

# construct an optimizer
params = [p for p in model_option_2.parameters() if p.requires_grad]
optimizer = torch.optim.SGD(params, lr=0.005,
                             momentum=0.9, weight_decay=0.0005)

# and a learning rate scheduler which decreases the learning rate by
# 10x every 3 epochs
lr_scheduler = torch.optim.lr_scheduler.StepLR(optimizer, step_size=3, gamma=0.
    ↪1)

for epoch in range(num_epochs):
    # train for one epoch, printing every 10 iterations
    train_one_epoch(model_option_2, optimizer, data_loader, device, epoch,
    ↪print_freq=10)
    # update the learning rate
    lr_scheduler.step()
    # evaluate on the test dataset
    evaluate(model_option_2, data_loader_test, device=device)

```

```

Epoch: [0] [ 0/60] eta: 0:00:52 lr: 0.000090 loss: 1.3941 (1.3941)
loss_classifier: 0.6496 (0.6496) loss_box_reg: 0.0400 (0.0400)
loss_objectness: 0.6887 (0.6887) loss_rpn_box_reg: 0.0158 (0.0158) time:
0.8816 data: 0.4675 max mem: 6172
Epoch: [0] [10/60] eta: 0:00:18 lr: 0.000936 loss: 1.3941 (1.3767)
loss_classifier: 0.6156 (0.5823) loss_box_reg: 0.0700 (0.0690)
loss_objectness: 0.6834 (0.6828) loss_rpn_box_reg: 0.0387 (0.0426) time:
0.3769 data: 0.0491 max mem: 6172
Epoch: [0] [20/60] eta: 0:00:14 lr: 0.001783 loss: 1.2109 (1.2129)
loss_classifier: 0.4132 (0.4440) loss_box_reg: 0.0804 (0.0921)
loss_objectness: 0.6480 (0.6397) loss_rpn_box_reg: 0.0345 (0.0371) time:
0.3255 data: 0.0079 max mem: 6172
Epoch: [0] [30/60] eta: 0:00:10 lr: 0.002629 loss: 0.9969 (1.1282)
loss_classifier: 0.2775 (0.3978) loss_box_reg: 0.1245 (0.1165)
loss_objectness: 0.5145 (0.5772) loss_rpn_box_reg: 0.0309 (0.0367) time:
0.3255 data: 0.0084 max mem: 6172
Epoch: [0] [40/60] eta: 0:00:06 lr: 0.003476 loss: 0.8696 (1.0501)
loss_classifier: 0.2537 (0.3638) loss_box_reg: 0.1726 (0.1344)
loss_objectness: 0.3614 (0.5140) loss_rpn_box_reg: 0.0379 (0.0378) time:
0.3286 data: 0.0092 max mem: 6172
Epoch: [0] [50/60] eta: 0:00:03 lr: 0.004323 loss: 0.7329 (0.9764)
loss_classifier: 0.2304 (0.3390) loss_box_reg: 0.1774 (0.1427)
loss_objectness: 0.2598 (0.4581) loss_rpn_box_reg: 0.0372 (0.0365) time:
0.3278 data: 0.0094 max mem: 6311
Epoch: [0] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.5542 (0.9140)
loss_classifier: 0.2203 (0.3174) loss_box_reg: 0.1513 (0.1441)
loss_objectness: 0.1963 (0.4166) loss_rpn_box_reg: 0.0310 (0.0359) time:
0.3243 data: 0.0087 max mem: 6311
Epoch: [0] Total time: 0:00:20 (0.3383 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:22 model_time: 0.0861 (0.0861) evaluator_time:
0.0043 (0.0043) time: 0.4502 data: 0.3587 max mem: 6311
Test: [49/50] eta: 0:00:00 model_time: 0.0360 (0.0385) evaluator_time:
0.0024 (0.0030) time: 0.0444 data: 0.0042 max mem: 6311
Test: Total time: 0:00:02 (0.0576 s / it)
Averaged stats: model_time: 0.0360 (0.0385) evaluator_time: 0.0024 (0.0030)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.045
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.150
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.006
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.001
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.098
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.047
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.255

```

```

Average Recall      (AR) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.352
Average Recall      (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall      (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.017
Average Recall      (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.386
Epoch: [1] [ 0/60]  eta: 0:00:48  lr: 0.005000  loss: 0.8100 (0.8100)
loss_classifier: 0.2839 (0.2839)  loss_box_reg: 0.2739 (0.2739)
loss_objectness: 0.1958 (0.1958)  loss_rpn_box_reg: 0.0563 (0.0563)  time:
0.8016  data: 0.3917  max mem: 6311
Epoch: [1] [10/60]  eta: 0:00:18  lr: 0.005000  loss: 0.4534 (0.5334)
loss_classifier: 0.1523 (0.1795)  loss_box_reg: 0.1344 (0.1588)
loss_objectness: 0.1527 (0.1623)  loss_rpn_box_reg: 0.0270 (0.0328)  time:
0.3706  data: 0.0435  max mem: 6311
Epoch: [1] [20/60]  eta: 0:00:14  lr: 0.005000  loss: 0.4769 (0.5590)
loss_classifier: 0.1566 (0.1898)  loss_box_reg: 0.1545 (0.1810)
loss_objectness: 0.1515 (0.1555)  loss_rpn_box_reg: 0.0293 (0.0327)  time:
0.3308  data: 0.0087  max mem: 6311
Epoch: [1] [30/60]  eta: 0:00:10  lr: 0.005000  loss: 0.5281 (0.5530)
loss_classifier: 0.1672 (0.1865)  loss_box_reg: 0.1838 (0.1859)
loss_objectness: 0.1397 (0.1485)  loss_rpn_box_reg: 0.0310 (0.0321)  time:
0.3331  data: 0.0094  max mem: 6311
Epoch: [1] [40/60]  eta: 0:00:06  lr: 0.005000  loss: 0.4674 (0.5419)
loss_classifier: 0.1397 (0.1805)  loss_box_reg: 0.1539 (0.1848)
loss_objectness: 0.1286 (0.1429)  loss_rpn_box_reg: 0.0317 (0.0337)  time:
0.3316  data: 0.0095  max mem: 6311
Epoch: [1] [50/60]  eta: 0:00:03  lr: 0.005000  loss: 0.4674 (0.5381)
loss_classifier: 0.1368 (0.1775)  loss_box_reg: 0.1667 (0.1909)
loss_objectness: 0.1212 (0.1365)  loss_rpn_box_reg: 0.0340 (0.0331)  time:
0.3398  data: 0.0091  max mem: 6618
Epoch: [1] [59/60]  eta: 0:00:00  lr: 0.005000  loss: 0.3847 (0.5144)
loss_classifier: 0.1274 (0.1689)  loss_box_reg: 0.1514 (0.1852)
loss_objectness: 0.0913 (0.1296)  loss_rpn_box_reg: 0.0162 (0.0307)  time:
0.3363  data: 0.0087  max mem: 6618
Epoch: [1] Total time: 0:00:20 (0.3444 s / it)
creating index...
index created!
Test: [ 0/50]  eta: 0:00:19  model_time: 0.0760 (0.0760)  evaluator_time:
0.0129 (0.0129)  time: 0.3836  data: 0.2936  max mem: 6618
Test: [49/50]  eta: 0:00:00  model_time: 0.0302 (0.0360)  evaluator_time:
0.0025 (0.0025)  time: 0.0410  data: 0.0047  max mem: 6618
Test: Total time: 0:00:02 (0.0546 s / it)
Averaged stats: model_time: 0.0302 (0.0360)  evaluator_time: 0.0025 (0.0025)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=  all | maxDets=100 ] = 0.203
Average Precision  (AP) @[ IoU=0.50      | area=  all | maxDets=100 ] = 0.591
Average Precision  (AP) @[ IoU=0.75      | area=  all | maxDets=100 ] = 0.046
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.004
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.224
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.125
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.364
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.382
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.017
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.419
Epoch: [2] [ 0/60] eta: 0:00:48 lr: 0.005000 loss: 0.2413 (0.2413)
loss_classifier: 0.0676 (0.0676) loss_box_reg: 0.0788 (0.0788)
loss_objectness: 0.0775 (0.0775) loss_rpn_box_reg: 0.0175 (0.0175) time:
0.8008 data: 0.4975 max mem: 6618
Epoch: [2] [10/60] eta: 0:00:18 lr: 0.005000 loss: 0.3504 (0.3711)
loss_classifier: 0.1112 (0.1168) loss_box_reg: 0.1132 (0.1311)
loss_objectness: 0.0987 (0.0956) loss_rpn_box_reg: 0.0214 (0.0276) time:
0.3653 data: 0.0497 max mem: 6618
Epoch: [2] [20/60] eta: 0:00:14 lr: 0.005000 loss: 0.3504 (0.3814)
loss_classifier: 0.1112 (0.1175) loss_box_reg: 0.1396 (0.1504)
loss_objectness: 0.0907 (0.0888) loss_rpn_box_reg: 0.0214 (0.0248) time:
0.3302 data: 0.0070 max mem: 6618
Epoch: [2] [30/60] eta: 0:00:10 lr: 0.005000 loss: 0.3664 (0.4233)
loss_classifier: 0.1116 (0.1308) loss_box_reg: 0.1791 (0.1742)
loss_objectness: 0.0747 (0.0918) loss_rpn_box_reg: 0.0218 (0.0265) time:
0.3367 data: 0.0092 max mem: 6618
Epoch: [2] [40/60] eta: 0:00:06 lr: 0.005000 loss: 0.4420 (0.4344)
loss_classifier: 0.1350 (0.1334) loss_box_reg: 0.2059 (0.1845)
loss_objectness: 0.0806 (0.0894) loss_rpn_box_reg: 0.0254 (0.0270) time:
0.3337 data: 0.0089 max mem: 6618
Epoch: [2] [50/60] eta: 0:00:03 lr: 0.005000 loss: 0.4420 (0.4320)
loss_classifier: 0.1350 (0.1326) loss_box_reg: 0.2158 (0.1873)
loss_objectness: 0.0700 (0.0858) loss_rpn_box_reg: 0.0254 (0.0263) time:
0.3382 data: 0.0089 max mem: 6618
Epoch: [2] [59/60] eta: 0:00:00 lr: 0.005000 loss: 0.3496 (0.4169)
loss_classifier: 0.1072 (0.1275) loss_box_reg: 0.1482 (0.1788)
loss_objectness: 0.0682 (0.0849) loss_rpn_box_reg: 0.0214 (0.0257) time:
0.3352 data: 0.0092 max mem: 6618
Epoch: [2] Total time: 0:00:20 (0.3442 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:19 model_time: 0.0727 (0.0727) evaluator_time:
0.0047 (0.0047) time: 0.3965 data: 0.3178 max mem: 6618
Test: [49/50] eta: 0:00:00 model_time: 0.0360 (0.0395) evaluator_time:
0.0024 (0.0027) time: 0.0450 data: 0.0048 max mem: 6618
Test: Total time: 0:00:02 (0.0584 s / it)
Averaged stats: model_time: 0.0360 (0.0395) evaluator_time: 0.0024 (0.0027)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox

```

```

Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.185
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.561
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.031
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.024
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.204
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.138
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.353
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.394
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.058
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.428
Epoch: [3] [ 0/60] eta: 0:00:43 lr: 0.000500 loss: 0.3568 (0.3568)
loss_classifier: 0.1135 (0.1135) loss_box_reg: 0.1811 (0.1811)
loss_objectness: 0.0437 (0.0437) loss_rpn_box_reg: 0.0185 (0.0185) time:
0.7236 data: 0.3614 max mem: 6618
Epoch: [3] [10/60] eta: 0:00:18 lr: 0.000500 loss: 0.2734 (0.3089)
loss_classifier: 0.0861 (0.0929) loss_box_reg: 0.1339 (0.1443)
loss_objectness: 0.0506 (0.0530) loss_rpn_box_reg: 0.0185 (0.0188) time:
0.3735 data: 0.0411 max mem: 6618
Epoch: [3] [20/60] eta: 0:00:14 lr: 0.000500 loss: 0.3210 (0.3442)
loss_classifier: 0.1048 (0.1073) loss_box_reg: 0.1563 (0.1613)
loss_objectness: 0.0506 (0.0559) loss_rpn_box_reg: 0.0164 (0.0197) time:
0.3404 data: 0.0088 max mem: 6618
Epoch: [3] [30/60] eta: 0:00:10 lr: 0.000500 loss: 0.3541 (0.3658)
loss_classifier: 0.1152 (0.1118) loss_box_reg: 0.1734 (0.1701)
loss_objectness: 0.0540 (0.0614) loss_rpn_box_reg: 0.0227 (0.0224) time:
0.3444 data: 0.0090 max mem: 6618
Epoch: [3] [40/60] eta: 0:00:07 lr: 0.000500 loss: 0.4011 (0.3683)
loss_classifier: 0.1207 (0.1127) loss_box_reg: 0.1689 (0.1714)
loss_objectness: 0.0627 (0.0623) loss_rpn_box_reg: 0.0211 (0.0219) time:
0.3413 data: 0.0092 max mem: 6660
Epoch: [3] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.3294 (0.3640)
loss_classifier: 0.1037 (0.1119) loss_box_reg: 0.1594 (0.1699)
loss_objectness: 0.0524 (0.0610) loss_rpn_box_reg: 0.0173 (0.0211) time:
0.3357 data: 0.0088 max mem: 6660
Epoch: [3] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.3350 (0.3582)
loss_classifier: 0.1049 (0.1099) loss_box_reg: 0.1551 (0.1673)
loss_objectness: 0.0474 (0.0597) loss_rpn_box_reg: 0.0190 (0.0213) time:
0.3360 data: 0.0083 max mem: 6660
Epoch: [3] Total time: 0:00:20 (0.3486 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:21 model_time: 0.1053 (0.1053) evaluator_time:
0.0199 (0.0199) time: 0.4358 data: 0.3027 max mem: 6660
Test: [49/50] eta: 0:00:00 model_time: 0.0353 (0.0391) evaluator_time:
0.0025 (0.0031) time: 0.0437 data: 0.0041 max mem: 6660
Test: Total time: 0:00:02 (0.0576 s / it)

```


Averaged stats: model_time: 0.0353 (0.0391) evaluator_time: 0.0025 (0.0031)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.257
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.654
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.145
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.014
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.284
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.150
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.448
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.457
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.033
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.500

Epoch: [4] [0/60] eta: 0:00:46 lr: 0.000500 loss: 0.2900 (0.2900)
loss_classifier: 0.0966 (0.0966) loss_box_reg: 0.1223 (0.1223)
loss_objectness: 0.0514 (0.0514) loss_rpn_box_reg: 0.0198 (0.0198) time:
0.7710 data: 0.3953 max mem: 6660
Epoch: [4] [10/60] eta: 0:00:18 lr: 0.000500 loss: 0.3474 (0.3662)
loss_classifier: 0.1049 (0.1149) loss_box_reg: 0.1795 (0.1757)
loss_objectness: 0.0507 (0.0555) loss_rpn_box_reg: 0.0190 (0.0201) time:
0.3713 data: 0.0431 max mem: 6660
Epoch: [4] [20/60] eta: 0:00:14 lr: 0.000500 loss: 0.3243 (0.3359)
loss_classifier: 0.0967 (0.1048) loss_box_reg: 0.1308 (0.1504)
loss_objectness: 0.0491 (0.0597) loss_rpn_box_reg: 0.0190 (0.0210) time:
0.3385 data: 0.0086 max mem: 6660
Epoch: [4] [30/60] eta: 0:00:10 lr: 0.000500 loss: 0.3399 (0.3623)
loss_classifier: 0.0969 (0.1114) loss_box_reg: 0.1323 (0.1696)
loss_objectness: 0.0490 (0.0590) loss_rpn_box_reg: 0.0211 (0.0224) time:
0.3450 data: 0.0097 max mem: 6660
Epoch: [4] [40/60] eta: 0:00:06 lr: 0.000500 loss: 0.2919 (0.3460)
loss_classifier: 0.0995 (0.1066) loss_box_reg: 0.1504 (0.1614)
loss_objectness: 0.0492 (0.0571) loss_rpn_box_reg: 0.0179 (0.0210) time:
0.3340 data: 0.0094 max mem: 6660
Epoch: [4] [50/60] eta: 0:00:03 lr: 0.000500 loss: 0.2919 (0.3454)
loss_classifier: 0.0964 (0.1062) loss_box_reg: 0.1393 (0.1623)
loss_objectness: 0.0477 (0.0556) loss_rpn_box_reg: 0.0170 (0.0214) time:
0.3386 data: 0.0087 max mem: 6660
Epoch: [4] [59/60] eta: 0:00:00 lr: 0.000500 loss: 0.3109 (0.3447)
loss_classifier: 0.1001 (0.1066) loss_box_reg: 0.1401 (0.1614)
loss_objectness: 0.0452 (0.0558) loss_rpn_box_reg: 0.0166 (0.0209) time:
0.3458 data: 0.0084 max mem: 6660
Epoch: [4] Total time: 0:00:21 (0.3501 s / it)
creating index...
index created!
Test: [0/50] eta: 0:00:20 model_time: 0.0852 (0.0852) evaluator_time:

```

0.0043 (0.0043)  time: 0.4017  data: 0.3109  max mem: 6660
Test: [49/50]  eta: 0:00:00  model_time: 0.0346 (0.0386)  evaluator_time:
0.0019 (0.0025)  time: 0.0436  data: 0.0046  max mem: 6660
Test: Total time: 0:00:02 (0.0572 s / it)
Averaged stats: model_time: 0.0346 (0.0386)  evaluator_time: 0.0019 (0.0025)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision  (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.251
Average Precision  (AP) @[ IoU=0.50      | area=   all | maxDets=100 ] = 0.694
Average Precision  (AP) @[ IoU=0.75      | area=   all | maxDets=100 ] = 0.099
Average Precision  (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision  (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.026
Average Precision  (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.279
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=  1 ] = 0.155
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 10 ] = 0.422
Average Recall     (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.443
Average Recall     (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall     (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.083
Average Recall     (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.480
Epoch: [5] [ 0/60]  eta: 0:00:49  lr: 0.000500  loss: 0.4341 (0.4341)
loss_classifier: 0.1329 (0.1329)  loss_box_reg: 0.2123 (0.2123)
loss_objectness: 0.0530 (0.0530)  loss_rpn_box_reg: 0.0359 (0.0359)  time:
0.8325  data: 0.4548  max mem: 6660
Epoch: [5] [10/60]  eta: 0:00:19  lr: 0.000500  loss: 0.3539 (0.3643)
loss_classifier: 0.1027 (0.1085)  loss_box_reg: 0.1730 (0.1745)
loss_objectness: 0.0511 (0.0566)  loss_rpn_box_reg: 0.0238 (0.0247)  time:
0.3892  data: 0.0481  max mem: 6660
Epoch: [5] [20/60]  eta: 0:00:14  lr: 0.000500  loss: 0.3401 (0.3573)
loss_classifier: 0.0937 (0.1069)  loss_box_reg: 0.1689 (0.1723)
loss_objectness: 0.0508 (0.0549)  loss_rpn_box_reg: 0.0208 (0.0232)  time:
0.3440  data: 0.0090  max mem: 6660
Epoch: [5] [30/60]  eta: 0:00:10  lr: 0.000500  loss: 0.2920 (0.3398)
loss_classifier: 0.0870 (0.1037)  loss_box_reg: 0.1556 (0.1642)
loss_objectness: 0.0412 (0.0504)  loss_rpn_box_reg: 0.0203 (0.0215)  time:
0.3412  data: 0.0103  max mem: 6660
Epoch: [5] [40/60]  eta: 0:00:07  lr: 0.000500  loss: 0.2920 (0.3483)
loss_classifier: 0.0958 (0.1059)  loss_box_reg: 0.1535 (0.1640)
loss_objectness: 0.0488 (0.0571)  loss_rpn_box_reg: 0.0180 (0.0214)  time:
0.3382  data: 0.0095  max mem: 6660
Epoch: [5] [50/60]  eta: 0:00:03  lr: 0.000500  loss: 0.3315 (0.3370)
loss_classifier: 0.0941 (0.1034)  loss_box_reg: 0.1555 (0.1590)
loss_objectness: 0.0458 (0.0543)  loss_rpn_box_reg: 0.0161 (0.0202)  time:
0.3399  data: 0.0088  max mem: 6660
Epoch: [5] [59/60]  eta: 0:00:00  lr: 0.000500  loss: 0.2708 (0.3371)
loss_classifier: 0.0890 (0.1039)  loss_box_reg: 0.1263 (0.1590)
loss_objectness: 0.0443 (0.0543)  loss_rpn_box_reg: 0.0126 (0.0198)  time:
0.3443  data: 0.0084  max mem: 6660

```

Epoch: [5] Total time: 0:00:21 (0.3539 s / it)
creating index...
index created!
Test: [0/50] eta: 0:00:22 model_time: 0.0900 (0.0900) evaluator_time: 0.0048 (0.0048) time: 0.4539 data: 0.3578 max mem: 6660
Test: [49/50] eta: 0:00:00 model_time: 0.0353 (0.0393) evaluator_time: 0.0023 (0.0024) time: 0.0437 data: 0.0047 max mem: 6660
Test: Total time: 0:00:02 (0.0582 s / it)
Averaged stats: model_time: 0.0353 (0.0393) evaluator_time: 0.0023 (0.0024)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox

Average Precision	(AP) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.275
Average Precision	(AP) @[IoU=0.50 area= all maxDets=100]	= 0.711
Average Precision	(AP) @[IoU=0.75 area= all maxDets=100]	= 0.133
Average Precision	(AP) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Precision	(AP) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.030
Average Precision	(AP) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.300
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 1]	= 0.149
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets= 10]	= 0.439
Average Recall	(AR) @[IoU=0.50:0.95 area= all maxDets=100]	= 0.451
Average Recall	(AR) @[IoU=0.50:0.95 area= small maxDets=100]	= -1.000
Average Recall	(AR) @[IoU=0.50:0.95 area=medium maxDets=100]	= 0.108
Average Recall	(AR) @[IoU=0.50:0.95 area= large maxDets=100]	= 0.486

Epoch: [6] [0/60] eta: 0:00:53 lr: 0.000050 loss: 0.1932 (0.1932)
loss_classifier: 0.0597 (0.0597) loss_box_reg: 0.0820 (0.0820)
loss_objectness: 0.0363 (0.0363) loss_rpn_box_reg: 0.0152 (0.0152) time: 0.8935 data: 0.5026 max mem: 6660
Epoch: [6] [10/60] eta: 0:00:19 lr: 0.000050 loss: 0.3186 (0.3075)
loss_classifier: 0.0996 (0.0959) loss_box_reg: 0.1537 (0.1470)
loss_objectness: 0.0481 (0.0471) loss_rpn_box_reg: 0.0187 (0.0175) time: 0.3885 data: 0.0519 max mem: 6660
Epoch: [6] [20/60] eta: 0:00:14 lr: 0.000050 loss: 0.3332 (0.3224)
loss_classifier: 0.1003 (0.0987) loss_box_reg: 0.1537 (0.1507)
loss_objectness: 0.0490 (0.0529) loss_rpn_box_reg: 0.0205 (0.0201) time: 0.3474 data: 0.0075 max mem: 6660
Epoch: [6] [30/60] eta: 0:00:10 lr: 0.000050 loss: 0.3388 (0.3358)
loss_classifier: 0.1008 (0.1022) loss_box_reg: 0.1621 (0.1597)
loss_objectness: 0.0469 (0.0532) loss_rpn_box_reg: 0.0224 (0.0207) time: 0.3506 data: 0.0083 max mem: 6660
Epoch: [6] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.3205 (0.3227)
loss_classifier: 0.0886 (0.0986) loss_box_reg: 0.1492 (0.1550)
loss_objectness: 0.0387 (0.0496) loss_rpn_box_reg: 0.0184 (0.0195) time: 0.3415 data: 0.0086 max mem: 6660
Epoch: [6] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.2788 (0.3159)
loss_classifier: 0.0868 (0.0960) loss_box_reg: 0.1301 (0.1506)
loss_objectness: 0.0387 (0.0496) loss_rpn_box_reg: 0.0165 (0.0197) time: 0.3415 data: 0.0094 max mem: 6660

Epoch: [6] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.3402 (0.3250)
 loss_classifier: 0.0984 (0.0998) loss_box_reg: 0.1544 (0.1547)
 loss_objectness: 0.0495 (0.0506) loss_rpn_box_reg: 0.0198 (0.0199) time:
 0.3459 data: 0.0089 max mem: 6660
 Epoch: [6] Total time: 0:00:21 (0.3570 s / it)
 creating index...
 index created!
 Test: [0/50] eta: 0:00:21 model_time: 0.0960 (0.0960) evaluator_time:
 0.0136 (0.0136) time: 0.4279 data: 0.3171 max mem: 6660
 Test: [49/50] eta: 0:00:00 model_time: 0.0355 (0.0396) evaluator_time:
 0.0023 (0.0026) time: 0.0436 data: 0.0044 max mem: 6660
 Test: Total time: 0:00:02 (0.0576 s / it)
 Averaged stats: model_time: 0.0355 (0.0396) evaluator_time: 0.0023 (0.0026)
 Accumulating evaluation results...
 DONE (t=0.02s).
 IoU metric: bbox
 Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.262
 Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.691
 Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.124
 Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = -1.000
 Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.020
 Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.294
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.177
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.439
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.451
 Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = -1.000
 Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.075
 Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.489
 Epoch: [7] [0/60] eta: 0:00:47 lr: 0.000050 loss: 0.2634 (0.2634)
 loss_classifier: 0.0798 (0.0798) loss_box_reg: 0.1262 (0.1262)
 loss_objectness: 0.0376 (0.0376) loss_rpn_box_reg: 0.0197 (0.0197) time:
 0.7959 data: 0.3951 max mem: 6660
 Epoch: [7] [10/60] eta: 0:00:19 lr: 0.000050 loss: 0.2942 (0.3072)
 loss_classifier: 0.0832 (0.0900) loss_box_reg: 0.1278 (0.1397)
 loss_objectness: 0.0528 (0.0606) loss_rpn_box_reg: 0.0178 (0.0169) time:
 0.3823 data: 0.0433 max mem: 6660
 Epoch: [7] [20/60] eta: 0:00:14 lr: 0.000050 loss: 0.2974 (0.3286)
 loss_classifier: 0.0893 (0.0997) loss_box_reg: 0.1457 (0.1539)
 loss_objectness: 0.0528 (0.0569) loss_rpn_box_reg: 0.0179 (0.0181) time:
 0.3416 data: 0.0083 max mem: 6660
 Epoch: [7] [30/60] eta: 0:00:10 lr: 0.000050 loss: 0.3806 (0.3508)
 loss_classifier: 0.1128 (0.1061) loss_box_reg: 0.1885 (0.1669)
 loss_objectness: 0.0503 (0.0581) loss_rpn_box_reg: 0.0206 (0.0198) time:
 0.3449 data: 0.0088 max mem: 6660
 Epoch: [7] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.3522 (0.3436)
 loss_classifier: 0.1106 (0.1044) loss_box_reg: 0.1847 (0.1639)
 loss_objectness: 0.0470 (0.0556) loss_rpn_box_reg: 0.0206 (0.0197) time:
 0.3469 data: 0.0088 max mem: 6660

Epoch: [7] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.2993 (0.3382)
 loss_classifier: 0.0824 (0.1029) loss_box_reg: 0.1278 (0.1615)
 loss_objectness: 0.0420 (0.0538) loss_rpn_box_reg: 0.0186 (0.0201) time:
 0.3434 data: 0.0093 max mem: 6660
 Epoch: [7] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.2993 (0.3301)
 loss_classifier: 0.0840 (0.1006) loss_box_reg: 0.1278 (0.1571)
 loss_objectness: 0.0422 (0.0531) loss_rpn_box_reg: 0.0151 (0.0193) time:
 0.3404 data: 0.0093 max mem: 6660
 Epoch: [7] Total time: 0:00:21 (0.3532 s / it)
 creating index...
 index created!
 Test: [0/50] eta: 0:00:20 model_time: 0.0931 (0.0931) evaluator_time:
 0.0146 (0.0146) time: 0.4030 data: 0.2940 max mem: 6660
 Test: [49/50] eta: 0:00:00 model_time: 0.0349 (0.0393) evaluator_time:
 0.0025 (0.0027) time: 0.0442 data: 0.0046 max mem: 6660
 Test: Total time: 0:00:02 (0.0576 s / it)
 Averaged stats: model_time: 0.0349 (0.0393) evaluator_time: 0.0025 (0.0027)
 Accumulating evaluation results...
 DONE (t=0.02s).
 IoU metric: bbox
 Average Precision (AP) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.256
 Average Precision (AP) @[IoU=0.50 | area= all | maxDets=100] = 0.664
 Average Precision (AP) @[IoU=0.75 | area= all | maxDets=100] = 0.102
 Average Precision (AP) @[IoU=0.50:0.95 | area= small | maxDets=100] = -1.000
 Average Precision (AP) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.049
 Average Precision (AP) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.282
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 1] = 0.134
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets= 10] = 0.445
 Average Recall (AR) @[IoU=0.50:0.95 | area= all | maxDets=100] = 0.450
 Average Recall (AR) @[IoU=0.50:0.95 | area= small | maxDets=100] = -1.000
 Average Recall (AR) @[IoU=0.50:0.95 | area=medium | maxDets=100] = 0.117
 Average Recall (AR) @[IoU=0.50:0.95 | area= large | maxDets=100] = 0.484
 Epoch: [8] [0/60] eta: 0:00:51 lr: 0.000050 loss: 0.3596 (0.3596)
 loss_classifier: 0.1115 (0.1115) loss_box_reg: 0.1839 (0.1839)
 loss_objectness: 0.0439 (0.0439) loss_rpn_box_reg: 0.0203 (0.0203) time:
 0.8504 data: 0.4487 max mem: 6660
 Epoch: [8] [10/60] eta: 0:00:19 lr: 0.000050 loss: 0.3550 (0.3759)
 loss_classifier: 0.1056 (0.1162) loss_box_reg: 0.1732 (0.1820)
 loss_objectness: 0.0556 (0.0556) loss_rpn_box_reg: 0.0219 (0.0221) time:
 0.3879 data: 0.0482 max mem: 6660
 Epoch: [8] [20/60] eta: 0:00:14 lr: 0.000050 loss: 0.3060 (0.3574)
 loss_classifier: 0.0911 (0.1085) loss_box_reg: 0.1659 (0.1757)
 loss_objectness: 0.0524 (0.0526) loss_rpn_box_reg: 0.0198 (0.0206) time:
 0.3401 data: 0.0086 max mem: 6660
 Epoch: [8] [30/60] eta: 0:00:10 lr: 0.000050 loss: 0.2957 (0.3427)
 loss_classifier: 0.0913 (0.1052) loss_box_reg: 0.1455 (0.1639)
 loss_objectness: 0.0428 (0.0545) loss_rpn_box_reg: 0.0152 (0.0190) time:
 0.3410 data: 0.0092 max mem: 6660

```

Epoch: [8] [40/60] eta: 0:00:07 lr: 0.000050 loss: 0.3075 (0.3406)
loss_classifier: 0.0925 (0.1054) loss_box_reg: 0.1455 (0.1631)
loss_objectness: 0.0410 (0.0533) loss_rpn_box_reg: 0.0176 (0.0188) time:
0.3468 data: 0.0094 max mem: 6660
Epoch: [8] [50/60] eta: 0:00:03 lr: 0.000050 loss: 0.3086 (0.3364)
loss_classifier: 0.0925 (0.1048) loss_box_reg: 0.1502 (0.1603)
loss_objectness: 0.0447 (0.0523) loss_rpn_box_reg: 0.0182 (0.0190) time:
0.3458 data: 0.0088 max mem: 6746
Epoch: [8] [59/60] eta: 0:00:00 lr: 0.000050 loss: 0.2822 (0.3279)
loss_classifier: 0.0922 (0.1024) loss_box_reg: 0.1314 (0.1558)
loss_objectness: 0.0379 (0.0508) loss_rpn_box_reg: 0.0154 (0.0188) time:
0.3419 data: 0.0080 max mem: 6746
Epoch: [8] Total time: 0:00:21 (0.3553 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:21 model_time: 0.1226 (0.1226) evaluator_time:
0.0048 (0.0048) time: 0.4221 data: 0.2935 max mem: 6746
Test: [49/50] eta: 0:00:00 model_time: 0.0344 (0.0397) evaluator_time:
0.0024 (0.0026) time: 0.0440 data: 0.0047 max mem: 6746
Test: Total time: 0:00:02 (0.0577 s / it)
Averaged stats: model_time: 0.0344 (0.0397) evaluator_time: 0.0024 (0.0026)
Accumulating evaluation results...
DONE (t=0.01s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.238
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.690
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.074
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.018
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.265
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.146
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.417
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.422
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.017
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.464
Epoch: [9] [ 0/60] eta: 0:00:48 lr: 0.000005 loss: 0.4472 (0.4472)
loss_classifier: 0.1144 (0.1144) loss_box_reg: 0.2432 (0.2432)
loss_objectness: 0.0601 (0.0601) loss_rpn_box_reg: 0.0296 (0.0296) time:
0.8097 data: 0.4448 max mem: 6746
Epoch: [9] [10/60] eta: 0:00:18 lr: 0.000005 loss: 0.3388 (0.3266)
loss_classifier: 0.0903 (0.0994) loss_box_reg: 0.1642 (0.1565)
loss_objectness: 0.0532 (0.0527) loss_rpn_box_reg: 0.0172 (0.0180) time:
0.3740 data: 0.0471 max mem: 6746
Epoch: [9] [20/60] eta: 0:00:14 lr: 0.000005 loss: 0.3066 (0.3390)
loss_classifier: 0.0853 (0.1042) loss_box_reg: 0.1563 (0.1630)
loss_objectness: 0.0438 (0.0530) loss_rpn_box_reg: 0.0156 (0.0189) time:
0.3384 data: 0.0086 max mem: 6746

```

```

Epoch: [9] [30/60] eta: 0:00:10 lr: 0.000005 loss: 0.2782 (0.3328)
loss_classifier: 0.0883 (0.1025) loss_box_reg: 0.1316 (0.1588)
loss_objectness: 0.0438 (0.0528) loss_rpn_box_reg: 0.0155 (0.0187) time:
0.3469 data: 0.0094 max mem: 6830
Epoch: [9] [40/60] eta: 0:00:07 lr: 0.000005 loss: 0.2940 (0.3329)
loss_classifier: 0.0883 (0.1012) loss_box_reg: 0.1345 (0.1582)
loss_objectness: 0.0459 (0.0538) loss_rpn_box_reg: 0.0180 (0.0197) time:
0.3472 data: 0.0090 max mem: 6830
Epoch: [9] [50/60] eta: 0:00:03 lr: 0.000005 loss: 0.3042 (0.3365)
loss_classifier: 0.0942 (0.1027) loss_box_reg: 0.1346 (0.1599)
loss_objectness: 0.0474 (0.0544) loss_rpn_box_reg: 0.0180 (0.0194) time:
0.3460 data: 0.0093 max mem: 6830
Epoch: [9] [59/60] eta: 0:00:00 lr: 0.000005 loss: 0.3316 (0.3402)
loss_classifier: 0.1007 (0.1041) loss_box_reg: 0.1413 (0.1616)
loss_objectness: 0.0547 (0.0547) loss_rpn_box_reg: 0.0188 (0.0199) time:
0.3421 data: 0.0087 max mem: 6830
Epoch: [9] Total time: 0:00:21 (0.3532 s / it)
creating index...
index created!
Test: [ 0/50] eta: 0:00:21 model_time: 0.0845 (0.0845) evaluator_time:
0.0053 (0.0053) time: 0.4263 data: 0.3352 max mem: 6830
Test: [49/50] eta: 0:00:00 model_time: 0.0336 (0.0394) evaluator_time:
0.0020 (0.0025) time: 0.0425 data: 0.0044 max mem: 6830
Test: Total time: 0:00:02 (0.0582 s / it)
Averaged stats: model_time: 0.0336 (0.0394) evaluator_time: 0.0020 (0.0025)
Accumulating evaluation results...
DONE (t=0.02s).
IoU metric: bbox
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.241
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.678
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.039
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.026
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.273
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.144
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.390
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.394
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = -1.000
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.042
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.430

```

After 10 epochs, we could notice that the option 1 (fine-tuned) model has lower triaining loss. And its AP and the AR metric are also better than the option 2 (different backbone) model.

Option 1 model AP and AR (bounding box):

Option 2 model AP and AR (bounding box):

Test option1 model and option2 model using Beatles picture

```
[56]: # First we download the picture
!wget https://upload.wikimedia.org/wikipedia/en/4/42/Beatles_-_Abbey_Road.jpg
Image.open('Beatles_-_Abbey_Road.jpg')
```

```
--2022-10-14 21:53:00--
https://upload.wikimedia.org/wikipedia/en/4/42/Beatles_-_Abbey_Road.jpg
Resolving upload.wikimedia.org (upload.wikimedia.org)... 103.102.166.240,
2001:df2:e500:ed1a::2:b
Connecting to upload.wikimedia.org
(upload.wikimedia.org)|103.102.166.240|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 89688 (88K) [image/jpeg]
Saving to: 'Beatles_-_Abbey_Road.jpg.2'

Beatles_-_Abbey_Roa 100%[=====] 87.59K --.-KB/s in 0.003s

2022-10-14 21:53:01 (26.6 MB/s) - 'Beatles_-_Abbey_Road.jpg.2' saved
[89688/89688]
```

[56]:



Validation of option 1 model:


```
[94]: from torchvision import transforms
test_Beatles_Image = Image.open('Beatles_-_Abbey_Road.jpg').convert("RGB")
test_Beatles_Image = test_Beatles_Image.resize((450,450))
test_Beatles_Image_Tensor = transforms.ToTensor()(test_Beatles_Image)

model.eval()
with torch.no_grad():
    predicted_Output = model([test_Beatles_Image_Tensor.to(device)])

predicted_Output
```

```
[94]: [{'boxes': tensor([[132.7701, 230.7768, 202.8403, 383.5952],
                        [215.5315, 238.9845, 302.2977, 396.4629],
                        [ 32.9992, 229.8836, 116.1947, 382.2405],
                        [311.6980, 236.8884, 400.8449, 411.6074],
                        [336.9757, 210.0914, 345.8099, 244.1632],
                        [ 78.9452, 225.4330, 193.7413, 382.2825],
                        [ 49.1275, 230.3554, 333.3038, 389.2376]]), device='cuda:0'),
      'labels': tensor([1, 1, 1, 1, 1, 1, 1], device='cuda:0'),
      'scores': tensor([0.9827, 0.9757, 0.9483, 0.8186, 0.1556, 0.0613, 0.0561],
                        device='cuda:0'),
      'masks': tensor([[[[0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          ...,
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.]],
                        [[0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          ...,
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.]],
                        [[0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          ...,
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.],
                          [0., 0., 0., ..., 0., 0., 0.]]],
```

```
...,
```

```
[[[0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   ...,  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.] ]],
```

```
[[[0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   ...,  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.] ]],
```

```
[[[0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   ...,  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.],  
   [0., 0., 0., ..., 0., 0., 0.] ]], device='cuda:0')}]
```


```
[95]: Image.fromarray(predicted_Output[0]['masks'][0, 0].mul(255).byte().cpu().  
      ↪ numpy())
```

```
[95]:
```




```
[96]: Image.fromarray(predicted_Output[0]['masks'][1, 0].mul(255).byte().cpu().  
      ↪numpy())
```

[96]:



```
[97]: Image.fromarray(predicted_Output[0]['masks'][2, 0].mul(255).byte().cpu().  
      ↪numpy())
```

[97]:



```
[98]: Image.fromarray(predicted_Output[0]['masks'][3, 0].mul(255).byte().cpu().  
      ↪numpy())
```

[98]:



Draw the bounding box

Here we choose the top-5 scored box

*This “plot_img_bbox” function is referred to : <https://www.kaggle.com/code/konstanter/fasterrcnn-pytorch-maskdetection>

```
[89]: import matplotlib.pyplot as plt
import matplotlib.patches as patches
# Function to visualize bounding boxes in the image

def plot_img_bbox(img, target):
    # plot the image and bboxes
    # Bounding boxes are defined as follows: x-min y-min width height
```


```

fig, a = plt.subplots(1,1)
fig.set_size_inches(5,5)
a.imshow(img)
bounding_Boxes = target['boxes'].cpu().numpy()
for i in range(5):
    box = bounding_Boxes[i]
    x, y, width, height = box[0], box[1], box[2]-box[0], box[3]-box[1]
    rect = patches.Rectangle((x, y), width, height, linewidth = 2,
    edgecolor = 'r', facecolor = 'none')

    # Draw the bounding box on top of the image
    a.add_patch(rect)
plt.show()

```

```
[93]: plot_img_bbox(test_Beatles_Image, predicted_Output[0])
```



Validation of option 2 model:

The top-5 scored bounding boxes are:

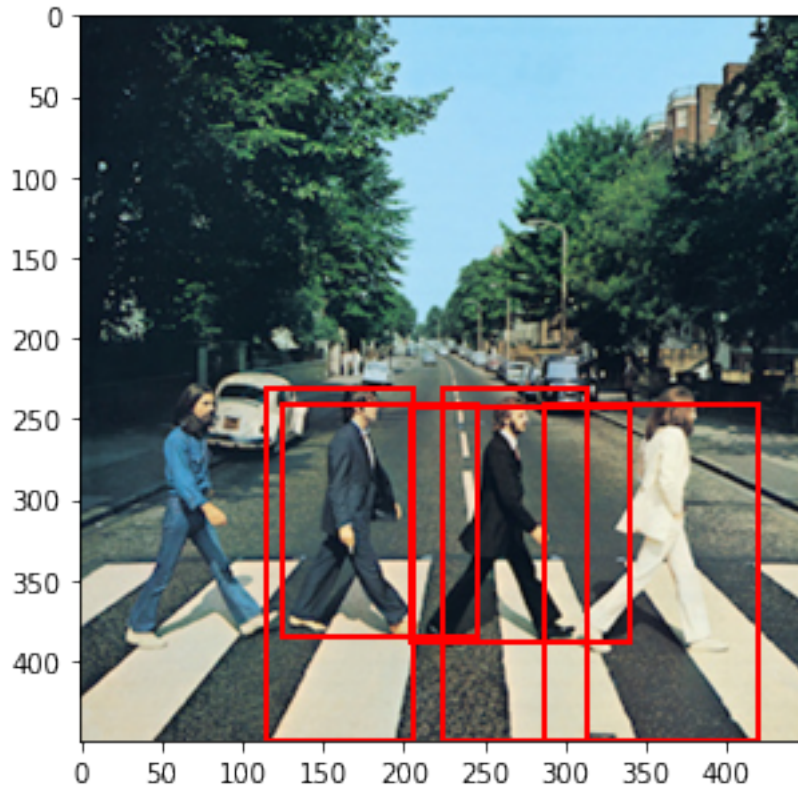
```
[90]: model_option_2.eval()
with torch.no_grad():
```

```

predicted_Output_2 = model_option_2([test_Beatles_Image_Tensor.
→to(device)])[0]

plot_img_bbox(test_Beatles_Image, predicted_Output_2)

```



So compared the top-5 scored bounding boxes of model1 and model2, we can notice that model 1 successfully select all 4 people occurred in the picture, while model 2 missed the most left people in the picture.

Thus, on this Beatles image, model 1 performs better than model 2.

```
[86]: predicted_Output_2
```

```
[86]: {'boxes': tensor([[223.6211, 231.1144, 313.8610, 450.0000],
                        [114.5353, 230.9814, 206.5275, 450.0000],
                        [203.9372, 243.5188, 340.9922, 389.3011],
                        [287.7523, 241.7243, 419.7959, 450.0000],
                        [124.6823, 240.9502, 246.3496, 385.5459],
                        [ 84.3286, 227.8522, 272.1978, 450.0000],
                        [195.9931, 214.4947, 435.7801, 450.0000],
                        [ 44.6954, 225.3032, 133.7435, 450.0000],
                        [ 60.3157, 211.2322, 185.7823, 450.0000],

```



```
      [ 22.0326, 157.3285, 114.2121, 414.2127]], device='cuda:0'),  
'labels': tensor([1, 1, 1, 1, 1, 1, 1, 1, 1, 1], device='cuda:0'),  
'scores': tensor([0.5101, 0.4976, 0.4890, 0.4545, 0.3267, 0.2958, 0.2167,  
0.2154, 0.0886,  
0.0630], device='cuda:0')}}}
```