

CS-GY 6083 A: Principles of Database Systems

Schema Refinement and Normal Forms

supplementary material:
“Database Management Systems” Sec. 19.1-19.6
class notes

Prof. Julia Stoyanovich
Computer Science and Engineering at Tandon
Center for Data Science
New York University

Motivating Example

Hourly_Employees (ssn, name, rating, wages, hours)

<u>ssn</u>	name	rating	wages	hours
123-22-3666	Ann	5	10	40
231-31-5368	Bob	5	10	30
131-24-3650	Cate	2	7	30
434-26-3751	Dan	2	7	32
612-67-4134	Eve	5	10	40

Observe: the value in the **wages** column is **determined** by the values in the **rating** column

Redundancy!

Why is this a problem?

The Evils of Redundancy

- **Redundancy** is at the root of several problems associated with relational schemas
 - **redundant storage**
 - **update anomalies**: if one copy with redundant info is updated and other copies are not
 - **insertion anomalies**: cannot insert a tuple for an employee unless we know the hourly wage corresponding to his rating
 - **deletion anomalies**: if we delete all tuples with rating 5, we won't know the hourly wage for that rating
- **Integrity constraints** (ICs) can be used to identify schemas with such problems and to suggest refinements
- Main refinement technique: **decomposition**

Decomposition

Hourly_Employees (ssn, name, rating, wages, hours)

<u>ssn</u>	name	rating	wages	hours
123-22-3666	Ann	5	10	40
231-31-5368	Bob	5	10	30
131-24-3650	Cate	2	7	30
434-26-3751	Dan	2	7	32
612-67-4134	Eve	5	10	40

Hourly_Employees2 (ssn, name, rating, hours) Rating_Wages(rating, wages)

<u>ssn</u>	name	rating	hours
123-22-3666	Ann	5	40
231-31-5368	Bob	5	30
131-24-3650	Cate	2	30
434-26-3751	Dan	2	32
612-67-4134	Eve	5	40

<u>rating</u>	wages
1	6
2	7
3	8
4	9
5	10

Functional Dependencies

A functional dependency is an integrity constraint (IC) that generalizes the concept of a key

Consider two non-empty sets of attributes of relation R :

$$A = \{A_1, A_2, \dots, A_n\} \quad \text{and} \quad B = \{B_1, B_2, \dots, B_m\}$$

A **functional dependency** (FD) $A \rightarrow B$ holds over relation R if, for every legal instance of R , and for any two tuples t_1 and t_2 , if t_1 and t_2 agree on the values of all attributes in **A** , then they agree on the values of all attributes in **B** .

(denoting sets of attributes with bold letters, e.g., **A**)

$$A_1 A_2 \dots A_n \rightarrow B_1 B_2 \dots B_m$$

Functional Dependencies

- An FD is a statement about **all legal instances** of R
 - An FD is given, **not** determined based on the data
 - Given a legal instance of R , we can check if it violates some FD, but we **cannot tell if an FD holds** over all legal instances

In particular, if K is a **candidate key** for R , then K **functionally determines** all other attributes in R .

Functional Dependencies: Example

Notation: refer to relation R by the first letters of its attribute names.

R: KABCD

K is a key

<u>K</u>	A	B	C	D
k1	a1	b1	c1	d1
k2	a1	b1	c1	d2
k3	a1	b2	c2	d1
k4	a1	b3	c3	d1

Is this FD satisfied: $K \rightarrow ABCD$ *yes, we were told K is a key*

Is this FD satisfied: $AB \rightarrow C$ *yes*

Is this FD satisfied: $A \rightarrow B$ *no, violated for tuples k3 and k4*

Is this FD satisfied: $AB \rightarrow B$ *yes, known as a trivial FD*

What is *Functional* about FDs?

$$A_1 A_2 \dots A_n \rightarrow B$$

is called a functional dependency because there is a function that takes a list of values, one for each attribute on the left, and produces a value for the attributes on the right.

Importantly, this function is computed in a specific way: by looking up the value of B in a relation.

Example:

Rating_Wages

<u>rating</u>	wages
1	6
2	7
3	8
4	9
5	10

FDs and Redundancy

- Given a schema, we need to decide whether it is well-designed (has no redundancy)
- If a schema does have redundancy, we may need to **decompose** it
- Decomposition is replacing relation R with several relations such that their sets of attributes together include all attributes R
- FDs help us reason about redundancy
 - Consider relation R with attributes ABC , denoted $R(A, B, C)$
 - If A is a candidate key and $A \rightarrow B$, there is **no redundancy**
 - If B is not a key and $B \rightarrow C$, then several tuples could have the same B value, so they'll have the same C value - **redundancy!**

Reasoning about FDs

Given a set of FDs that relation R satisfies, infer additional FDs that hold in R .

Two sets of FDs \mathbf{S} and \mathbf{T} are **equivalent** if the set of instances satisfying \mathbf{S} is the same as the set of instances satisfying \mathbf{T} .

A set of FDs \mathbf{S} **follows** from a set of FDs \mathbf{T} if every relation instance that satisfies all the FDs in \mathbf{T} also satisfies all the FDs in \mathbf{S} .

Basic Rules

The splitting rule

$A_1 A_2 \dots A_n \rightarrow B_1 B_2 \dots B_m$ can be replaced by a set of m rules

$A_1 A_2 \dots A_n \rightarrow B_i$ with $i = 1, 2, \dots, m$

The combining rule

$A_1 A_2 \dots A_n \rightarrow B_i$ with $i = 1, 2, \dots, m$ can be replaced by

$A_1 A_2 \dots A_n \rightarrow B_1 B_2 \dots B_m$

The transitive rule

$A_1, A_2, \dots, A_n \rightarrow B_1, B_2, \dots, B_m \wedge$
 $B_1, B_2, \dots, B_m \rightarrow C_1, C_2, \dots, C_k \Rightarrow$
 $A_1, A_2, \dots, A_n \rightarrow C_1, C_2, \dots, C_k$

Proof for

$A \rightarrow B \wedge B \rightarrow C \Rightarrow A \rightarrow C$

Proof of the transitive rule

$$R(A, B, C) \quad A \rightarrow B \wedge B \rightarrow C \Rightarrow A \rightarrow C$$

- Consider tuples $t_1(a, b_1, c_1), t_2(a, b_2, c_2)$
- Since $A \rightarrow B$ and $t_1.A = t_2.A = a$, then it must be the case that $t_1.B = t_2.B$, i.e., that $b_1 = b_2$
- Similarly, since $B \rightarrow C$, and since $t_1.B = t_2.B$ then $t_1.C = t_2.C$, i.e., $c_1 = c_2$

Closure of a set of attributes

Suppose $\mathbf{A} = \{A_1, \dots, A_n\}$ is a set of attributes and \mathbf{S} is a set of FDs.

The *closure* of \mathbf{A} under the FDs in \mathbf{S} is the set of attributes \mathbf{B} s.t. every relation that satisfies all the FDs in \mathbf{S} also satisfies $A \rightarrow B$

We denote the closure of $\{A_1, A_2, \dots, A_n\}$ by $\{A_1, A_2, \dots, A_n\}^+$

Note that $\{A_1, A_2, \dots, A_n\} \subseteq \{A_1, A_2, \dots, A_n\}^+$ Why?

Computing the closure of a set of attributes

Algorithm *AttributeClosure*

Input: a set of attributes $\{A_1, A_2, \dots, A_n\}$ and a set of FDs \mathbf{S}

Output: the closure $\{A_1, A_2, \dots, A_n\}^+$

1. Split the FDs of \mathbf{S} using the splitting rule, so that each FD has one attribute on the right
2. Initialize $\{A_1, A_2, \dots, A_n\}^+ \leftarrow \{A_1, A_2, \dots, A_n\}$
3. Repeatedly search for some FD $B_1, B_2, \dots, B_m \rightarrow C$ such that
 $\{B_1, B_2, \dots, B_m\} \subseteq \{A_1, A_2, \dots, A_n\}^+ \wedge C \notin \{A_1, A_2, \dots, A_n\}^+$
4. Stop when no more attributes can be added to $\{A_1, A_2, \dots, A_n\}^+$

Example

$$R(A, B, C, D) \quad S = \{AB \rightarrow C, C \rightarrow D, D \rightarrow A\}$$

$$\{A, B, D\}^+ = \{A, B, D\} \cup_{AB \rightarrow C} \{C\} = \{A, B, C, D\}$$

$$\{C, D\}^+ = \{C, D\} \cup_{D \rightarrow A} \{A\} = \{A, C, D\}$$

Note that B is not in the closure of {CD}. Generally, since B does not appear on the right-hand side of any rule, it will not be in the closure of any set that does not already contain B.

Example

$$R(A, B, C, D) \quad S = \{A \rightarrow B, BC \rightarrow A\}$$

$$\{A\}^+ =$$

$$\{B\}^+ =$$

$$\{B, C\}^+ =$$

$$\{A, B, C\}^+ =$$

$$\{B, C, D\}^+ =$$

Example (solution)

$$R(A, B, C, D) \quad S = \{A \rightarrow B, BC \rightarrow A\}$$

$$\{A\}^+ =$$

$$\{B\}^+ =$$

$$\{B, C\}^+ =$$

$$\{A, B, C\}^+ =$$

$$\{B, C, D\}^+ =$$

What can we do with the closure?

Given \mathbf{S} and an FD not in \mathbf{S} , we can compute whether that FD follows from \mathbf{S} .

Consider $B_1B_2 \dots B_m \rightarrow C$

1. Compute the closure of the attributes on the left, $\{B_1, B_2, \dots, B_m\}^+$
2. Check whether the attribute on the right, C , is in the closure.
If so - FD follows from \mathbf{S} , otherwise it does not.

More generally, $A_1A_2 \dots A_n \rightarrow B_1B_2 \dots B_m$

follows from \mathbf{S} if and only if $\{B_1, B_2, \dots, B_m\} \subseteq \{A_1, A_2, \dots, A_n\}^+$

Example

$$R(A, B, C, D) \quad S = \{AB \rightarrow C, C \rightarrow D, D \rightarrow A\}$$

Check whether the following FDs follow from **S**:

Compute the closure of the left-hand side, check whether the attribute on the right-hand side is in the closure. If so - yes, if not - no.

$$C \rightarrow A \quad \{C\}^+ = \{A, C, D\} \quad \text{yes}$$

$$CD \rightarrow B \quad \{C, D\}^+ = \{A, C, D\} \quad \text{no}$$

$$AB \rightarrow D \quad \{A, B\}^+ = \{A, B, C, D\} \quad \text{yes}$$

Correctness of the closure algorithm

The algorithm is **sound**: it computes **only** the true FDs

The algorithm is **complete**: it computes **all** the true FDs

In a relation R , for which set or sets of attributes does the closure correspond to all attributes of R ?

Closure of a Set of FDs

- For a given FD, we can decide whether it follows from a given set of FDs **S** using the closure of the set of attributes algorithm
- Alternatively, we can compute the closure of the set of FDs **S** and check whether the FD in question is in that set. We can do this using **Armstrong's axioms**.

Reflexivity
(trivial FDs)

*If $\{B_1, B_2, \dots, B_m\} \subseteq \{A_1, A_2, \dots, A_n\}$
then $A_1 A_2 \dots A_n \rightarrow B_1 B_2 \dots B_m$*

Augmentation

*If $A_1 A_2 \dots A_n \rightarrow B_1 B_2 \dots B_m$
then $A_1 A_2 \dots A_n C_1 C_2 \dots C_k \rightarrow B_1 B_2 \dots B_m C_1 C_2 \dots C_k$*

Transitivity

*If $A_1 A_2 \dots A_n \rightarrow B_1 B_2 \dots B_m$ and $B_1 B_2 \dots B_m \rightarrow C_1 C_2 \dots C_k$
then $A_1 A_2 \dots A_n \rightarrow C_1 C_2 \dots C_k$*

Minimal basis of a set of FDs

- For a given relation R , there may exist several sets of FDs that are equivalent:
 - they give rise to the same closures of all subsets of R 's attributes
 - the same sets of FDs follow from them
 - all such equivalent sets of FDs are called *bases for S in R*
- A *minimal basis B* is a set of FDs that satisfies 3 conditions
 1. All FDs in B have 1 attribute on the right (are in a standard form)
 2. If any FD is removed from B , the result is no longer a basis
 3. If for any FD in B we remove 1 attribute on the left, the result is no longer a basis

Example

$$R(A,B,C) \quad S = \{A \rightarrow BC, B \rightarrow AC, C \rightarrow AB\}$$

What is the full set of FDs with 1 attribute on the right that follow from **S**?

To work with **S** more easily, use the splitting rule to transform each FD into 2 FDs

Is **S** a minimal basis? $S = \{A \rightarrow B, A \rightarrow C, B \rightarrow A, B \rightarrow C, C \rightarrow A, C \rightarrow B\}$

No, e.g., removing $A \rightarrow B$ does not change the closure of any attribute.

Give 2 different minimal bases for **S** in **R**.

$$S_1 = \{A \rightarrow B, B \rightarrow C, C \rightarrow A\}$$

$$S_2 = \{A \rightarrow C, C \rightarrow B, B \rightarrow A\}$$

Computing the minimal basis

Called “Minimal Cover” in Ramakrishnan & Gehrke, Section 19.6.2

Algorithm *MinimalBasis*

Input: Relation R , a set of FDs \mathbf{S} that hold in R .

Output: The set of FDs \mathbf{T} that forms a minimal basis for \mathbf{S} in R .

1. Transform the FDs in \mathbf{S} into a standard form: if an FD has one attribute on the right, add it to \mathbf{T} . Otherwise, if it has m attributes on the right, break it up into m FDs using the splitting rule (see slide 11) and add each to \mathbf{T} .
2. Compute the closure of each set of attributes in the powerset of R (except the empty set) under \mathbf{T} .
3. Minimize the left side of each FD in \mathbf{T} : For each FD in \mathbf{T} with 2 or more attributes on the left, check if an attribute can be deleted while preserving the closures computed in step 2.
4. Delete redundant FDs: check each FD in \mathbf{T} to see if it can be deleted while preserving the closures computed in step 2.

Example

$$R(A, B, C, D, E) \quad S = \{ABCD \rightarrow E, E \rightarrow D, A \rightarrow B, AC \rightarrow D\}$$

$$\text{Step 1: } T = \{ABCD \rightarrow E, E \rightarrow D, A \rightarrow B, AC \rightarrow D\}$$

$$\begin{array}{ll} \text{Step 2: } \{A\}^+ = \{A, B\} = \{A, B\}^+ & A \rightarrow B \\ & E \rightarrow D \\ \{E\}^+ = \{E, D\} & \\ \{A, C\}^+ = \{A, B, C, D, E\} & AC \rightarrow B, AC \rightarrow D, AC \rightarrow E \\ \{A, E\}^+ = \{A, B, D, E\} & AE \rightarrow B, AE \rightarrow D \end{array}$$

$$\text{Step 3: } T = \{AC \rightarrow E, E \rightarrow D, A \rightarrow B, AC \rightarrow D\}$$

$$\text{Step 4: } T = \{AC \rightarrow E, E \rightarrow D, A \rightarrow B\}$$

could we have dropped the first two FDs instead?

Closures and keys

Q: How can we tell if a set of attributes $A_1A_2 \dots A_n$ is a key of a relation R ?

A: If $\{A_1, A_2 \dots A_n\}^+ = \text{all the attributes in } R$

Q: How can we compute the keys for R ?

A: Find all sets of attributes that functionally determine all other attributes. *Is this enough?*

A: Make sure these sets are minimal!

Example

Movies(title, year, studio, president, address)

$FDs = \{TY \rightarrow S, S \rightarrow P, S \rightarrow A\}$

Compute the keys of *Movies*

$$\{S\}^+ = \{S, P, A\}$$

$$\{T\}^+ = \{T\}$$

$$\{Y\}^+ = \{Y\}$$

$$\{P\}^+ = \{P\}$$

$$\{A\}^+ = \{A\}$$

$$\{TY\}^+ = \{T, Y, S, P, A\}$$

Remember that there can be multiple candidate keys,
we just happen to have 1 candidate key here!

Example

Compute the keys of R

$$R(A, B, C, D) \quad S = \{A \rightarrow B, B \rightarrow C, C \rightarrow D\}$$

$$R(A, B, C, D, E) \quad S = \{AB \rightarrow C, C \rightarrow B, A \rightarrow D\}$$

Example (solution)

Compute the keys of R

$$R(A, B, C, D) \quad S = \{A \rightarrow B, B \rightarrow C, C \rightarrow D\}$$

$$R(A, B, C, D, E) \quad S = \{AB \rightarrow C, C \rightarrow B, A \rightarrow D\}$$

Computing a projection of a set of FDs

Suppose relation R is given, with its corresponding set of FDs \mathbf{S} . If we take a projection of R onto a set of attributes L , what can we say about the FDs of $\pi_L(R)$?

Algorithm *ProjectFDs*

Input: Relations R and $R1 = \pi_L(R)$. A set of FDs \mathbf{S} that hold in R .

Output: The set of FDs \mathbf{T} that hold in $R1$.

1. Compute the closure of each subset of attributes of $R1$ in \mathbf{S} , store the result in \mathbf{T} . Add to \mathbf{T} all non-trivial FDs $X \rightarrow A$ s.t. A is both in X^+ and an attribute of $R1$.
2. Remove from \mathbf{T} all FDs that involve attributes not in L (on either side).
3. Optionally compute the minimal basis of \mathbf{T} , remove FDs from \mathbf{T} that do not belong to the minimal basis.

Example

Compute a projection of the set of FDs when $R(ABCD)$ is projected onto ACD .

$R(ABCD) \quad A \rightarrow B ; B \rightarrow C ; C \rightarrow D$

$\pi_{ACD}(R)$

Compute closures of all subsets of attributes in the projected relation.

$$\{A\}^+ = \{A, B, C, D\}$$

$$\{C\}^+ = \{C, D\} = \{C, D\}^+$$

$$\{D\}^+ = \{D\}$$

we stop here, since any set that includes A will have the same closure as A alone

Compute FDs from these closures that involve only A, C, D on either side.

$$T = \{A \rightarrow C, A \rightarrow D, C \rightarrow D\}$$

Optionally remove redundant FDs, keeping only the minimal basis.

$$T = \{A \rightarrow C, C \rightarrow D\}$$

Why not simply take S and project each FD?

Recap so far

- What are functional dependencies (FDs)?
- We care about FDs because they allow us to reason about redundancy in a relation
- We can use FDs to compute the closure of a set of attributes
- We can use FDs to check which FDs are inferred from them
- We know how to project a set of FDs when the relation is projected onto a subset of its attributes

Next: using FDs to improve schema design (aka normalization)

Normal Forms

We consider the FDs that hold in a schema, and, based on that information, classify the schema as being in a certain **normal form**

- First normal form (1NF): all attributes have atomic values; the relational model guarantees 1NF
- Second normal form (2NF), of historical interest only, a less restrictive version of 3NF
- Third normal form (3NF)*
- Boyce-Codd normal form (BCNF)** - the Holy Grail

If R is in BCNF it is also in 3NF; if R is in 3NF it is also in 2NF.

All relations are in 1NF, by definition of the relational model. And so “all relations that are in 2NF are also in 1NF” is a truism. All relations are in 1NF, period.

Boyce-Codd Normal Form (BCNF)



Let R be a relation schema, S be the set of FDs given to hold over R .

R is in BCNF if, for every FD $A_1A_2 \dots A_n \rightarrow B_1B_2 \dots B_m$

one of the following statements is true:

1. The FD is **trivial**: $\{B_1, B_2, \dots, B_m\} \subseteq \{A_1, A_2, \dots, A_n\}$
2. $A_1A_2 \dots A_n$ is a candidate key of R
3. $A_1A_2 \dots A_n$ is a superkey of R

In a BCNF relation, the only set of attributes that determines values for other attributes is a key!

Example

Hourly_Employees (ssn, name, rating, wages, hours)

<u>ssn</u>	name	rating	wages	hours
123-22-3666	Ann	5	10	40
231-31-5368	Bob	5	10	30
131-24-3650	Cate	2	7	30
434-26-3751	Dan	2	7	32
612-67-4134	Eve	5	10	40

Is Hourly_Employees in BCNF?

FDs:

$S \rightarrow NRW$

$R \rightarrow W$

no

FDs:

$S \rightarrow NRW$

$SR \rightarrow W$

yes

FDs:

$S \rightarrow NRW$

$NR \rightarrow N$

yes

A relation with 2 attributes

R(A,B) e.g., SSN and Id are both candidate keys

ssn	id
123-22-3666	1
231-31-5368	2
131-24-3650	3
434-26-3751	4
612-67-4134	5

$$A \rightarrow B$$

$$B \rightarrow A$$

Is this relation in BCNF?

more generally, all 2-attribute relations are in BCNF

To see this, consider all cases of FDs that can hold, see that each case gives rise to a BCNF relation.

Decomposition of Relation Schema

A **decomposition** of relation R consists of replacing R by two or more relations such that:

- each new relation contains a subset of the attributes of R (and no new attributes), and
- each attribute of R appears in at least one of the new relations.

Decomposition

Hourly_Employees : SNRWH

<u>ssn</u>	name	rating	wages	hours
123-22-3666	Ann	5	10	40
231-31-5368	Bob	5	10	30
131-24-3650	Cate	2	7	30
434-26-3751	Dan	2	7	32
612-67-4134	Eve	5	10	40

FDs:

$S \rightarrow NRWH$

$R \rightarrow W$

Hourly_Employees2 : SNRH

<u>ssn</u>	name	rating	hours
123-22-3666	Ann	5	40
231-31-5368	Bob	5	30
131-24-3650	Cate	2	30
434-26-3751	Dan	2	32
612-67-4134	Eve	5	40

Rating_Wages : RW

<u>rating</u>	wages
1	6
2	7
3	8
4	9
5	10

Decomposition into BCNF

Let R be a relation schema, S be the set of FDs given to hold over R . We decompose R by considering FDs that violate BCNF.

Suppose that the following non-trivial FD violates BCNF

$$A_1 A_2 \dots A_n \rightarrow B$$

This means that

1. $A_1 A_2 \dots A_n$ functionally determine B
2. $A_1 A_2 \dots A_n$ is **not** a key

Intuition

1. Use the offending FD to make a BCNF relation (one where the left side is a candidate key

What relation is that? $A_1 A_2 \dots A_n \rightarrow ?$

2. Create another relation from which redundancy due to the offending FD has been removed

Decomposition into BCNF

Let R be a relation schema, S be the set of FDs given to hold over R . We decompose R by considering FDs that violate BCNF.

Algorithm *BCNFDecomposition*

Input: Relation R , a set of FDs S that hold in R .

Output: A decomposition of R into a set of relations, all of which are in BCNF.

1. Check whether R is in BCNF. If so, return $\{R\}$.
2. Otherwise, let $A_1A_2 \dots A_n \rightarrow B$ be an FD that violates BCNF.
 - 2.1. Use *AttributeClosure* to compute $\{A_1, A_2, \dots, A_n\}^+$
 - 2.2. Decompose R into $R_1 = R - B$ and $R_2 = \{A_1A_2 \dots A_n\}^+$
 - 2.3. Use *ProjectFDs* to compute FDs of R_1 and R_2
 - 2.4. Recursively decompose R_1 and R_2 using *BCNFDecomposition*


Decomposition into BCNF

Example Hourly_Employees : SNRWH

$\{S \rightarrow N, S \rightarrow R, S \rightarrow W, S \rightarrow H, R \rightarrow W\}$

$R \rightarrow W$ violates BCNF

decompose into SNRH and RW



$\{S \rightarrow N, S \rightarrow R, S \rightarrow H\}$ $R \rightarrow W$ FDs on the decomposed relations

done!

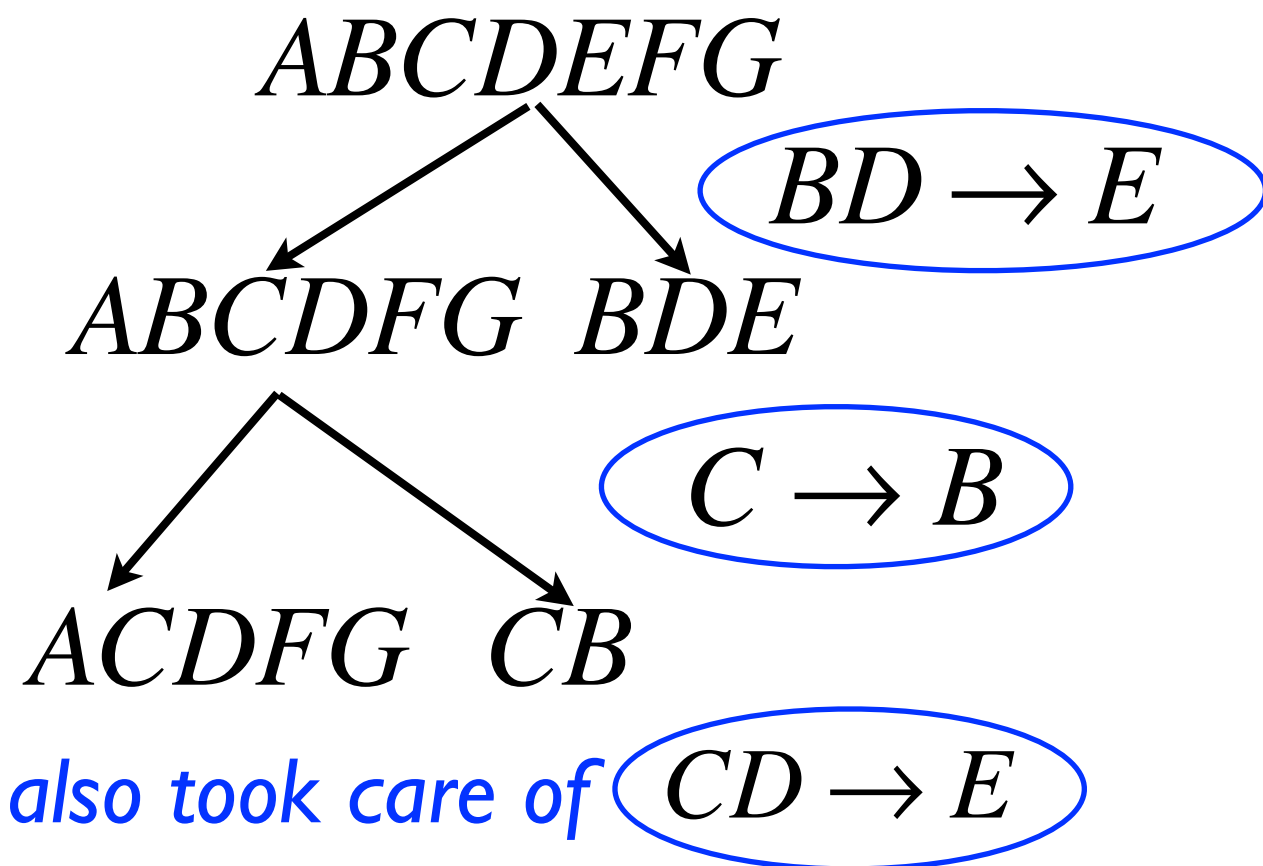
Decomposition into BCNF

If multiple FDs violate BCNF, different orders of decomposition are possible, leading to different results

Example: R is ABCDEFG, key A

violate BCNF

FDs: $A \rightarrow BCDEFG$ $C \rightarrow B$ $BD \rightarrow E$ $CD \rightarrow E$



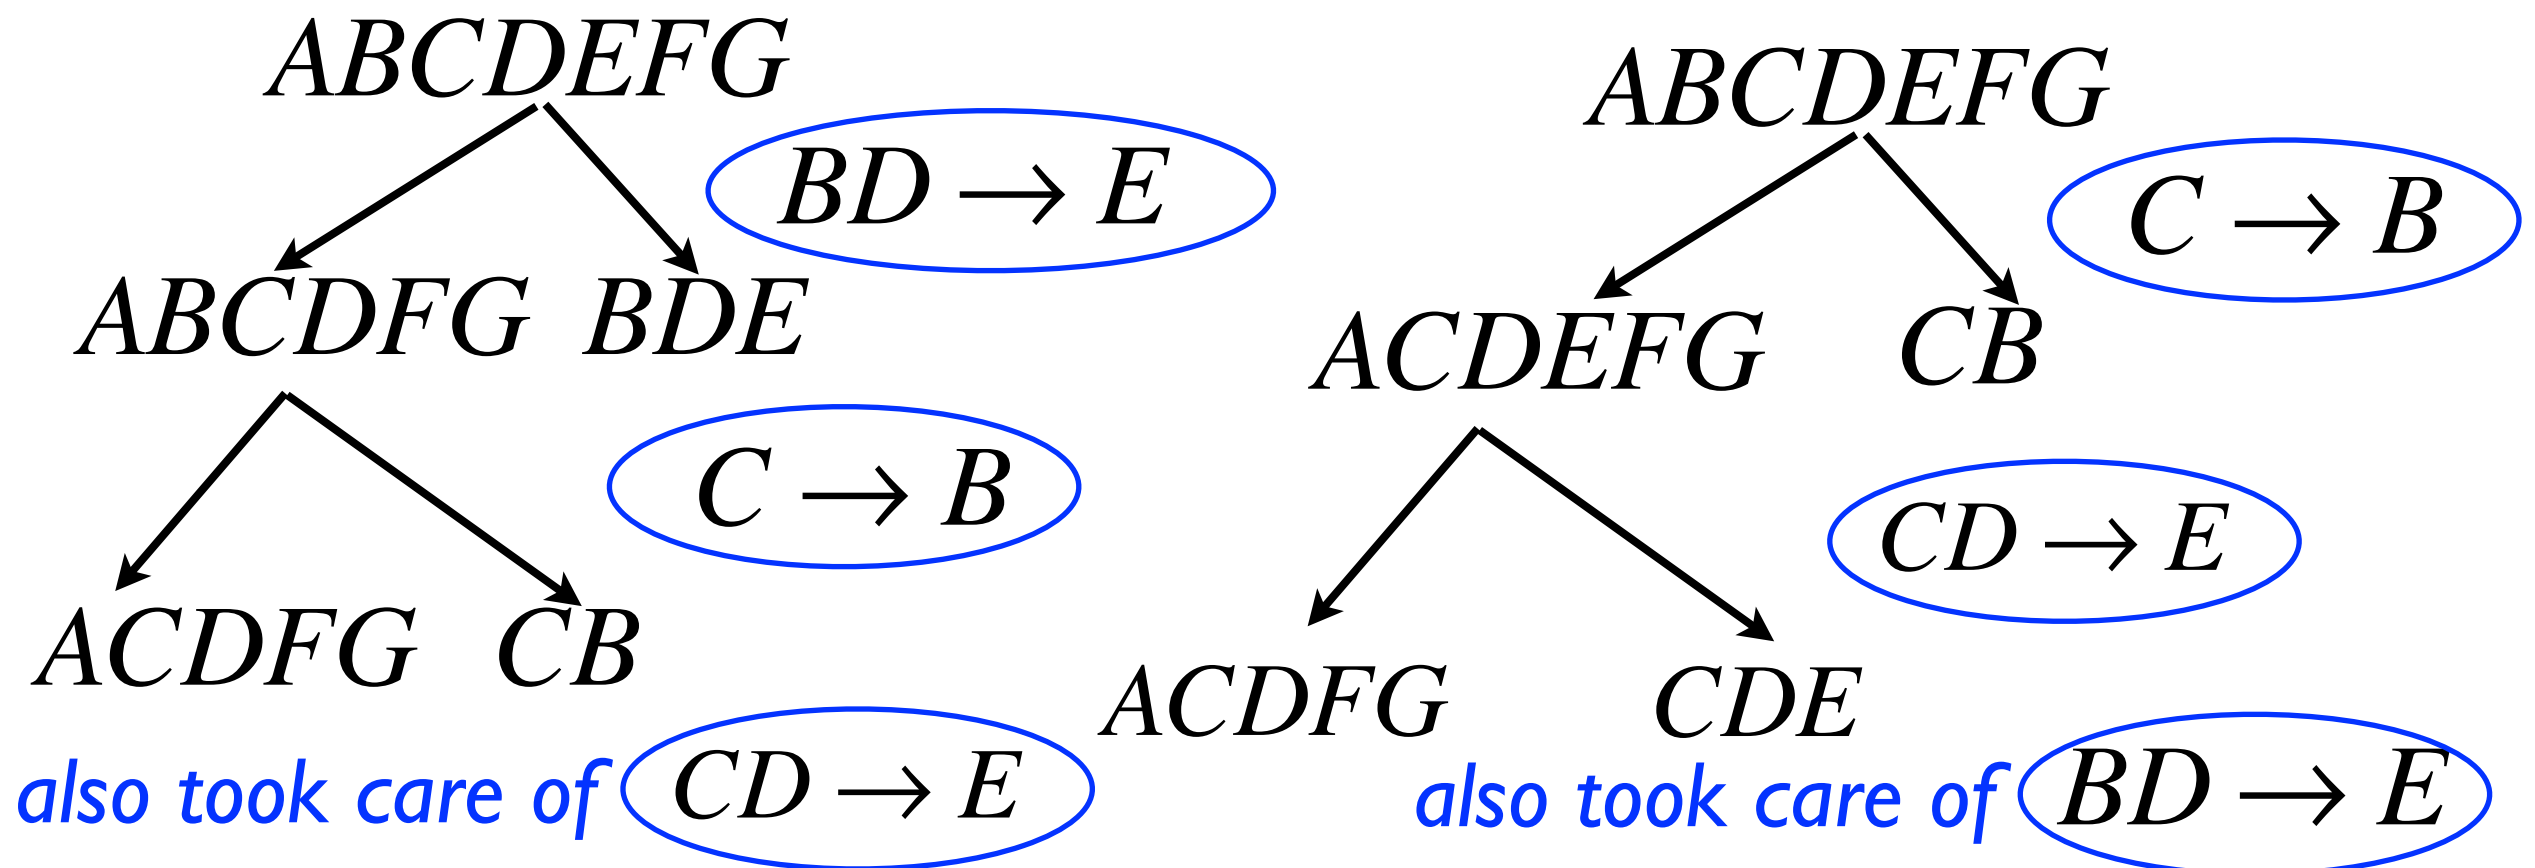
Decomposition into BCNF

If multiple FDs violate BCNF, different orders of decomposition are possible, leading to different results

Example: R is ABCDEFG, key A

violate BCNF

FDs: $A \rightarrow BCDEFG$ $C \rightarrow B$ $BD \rightarrow E$ $CD \rightarrow E$



Properties of a decomposition

- If a relation was not in BCNF, and we decomposed it into a set of BCNF relations, we **eliminate the anomalies due to redundancy**
- A good decomposition should additionally have two properties:
 1. **Recoverability of information**: Can we recover the original relation from the tuples of its decomposition (by joining)? If so, we have a **lossless join** decomposition.
 2. **Preservation of dependencies**: When we reconstruct the original relation, does it still satisfy all the original FDs? If so, we have a **dependency-preserving** decomposition.

Lossless join decomposition

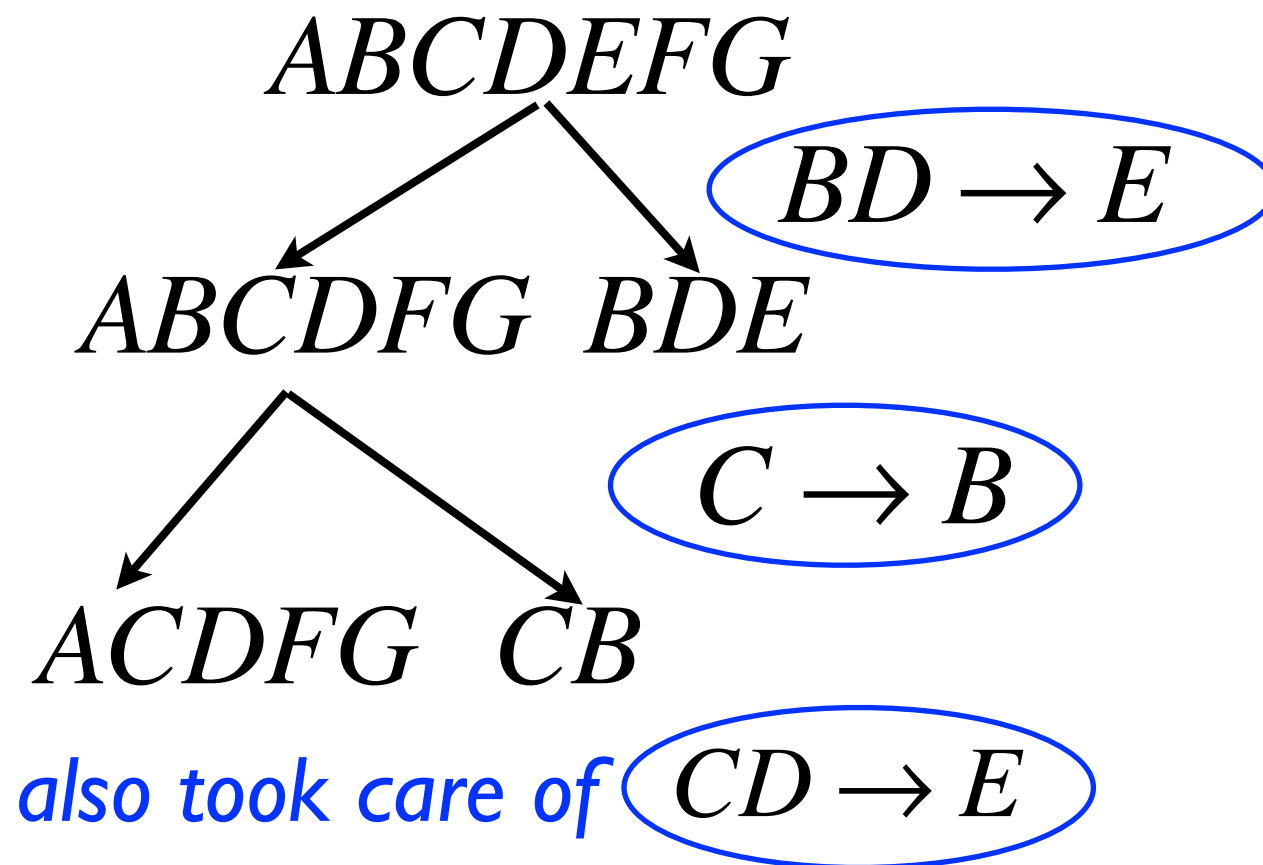
Any decomposition using *BCNFDecomposition* will result in lossless join

- This is because we decompose using an FD, and can reconstruct the original relation (schema and tuples), by taking a natural join. This will compute all the original tuples and no additional ones.
- If we decomposed without an FD - we may generate more tuples by the natural join than were there originally

Recall: Decomposition into BCNF

R: $ABCDEFG$ FDs: $A \rightarrow BCDEFG$ *violate BCNF*

$C \rightarrow B$ $BD \rightarrow E$ $CD \rightarrow E$



Are all of the original FDs still enforced?

No, this decomposition is not dependency-preserving!

Dependency-preserving decomposition

In a **dependency-preserving decomposition**, each original dependency can be checked by looking at just 1 table in the result of the decomposition

R: $ABCDEFG$ FDs: $C \rightarrow B$ $BD \rightarrow E$ $CD \rightarrow E$

$R1(ACDFG)$ $R2(CB)$ $R3(BDE)$ not dependency-preserving
 $CD \rightarrow E$ not enforced

$R1(ACDFG)$ $R2(CB)$ $R3(CDE)$ not dependency-preserving
 $BD \rightarrow E$ not enforced

A dependency-preserving decomposition into BCNF is
not guaranteed to exist.

Example

$R(ABCD)$

$A \rightarrow B \quad B \rightarrow D \quad AD \rightarrow C \quad BC \rightarrow A$

Decompose R into BCNF. Show keys, projected FDs.

Example (solution)

$R(ABCD)$

$A \rightarrow B \quad B \rightarrow D \quad AD \rightarrow C \quad BC \rightarrow A$

Decompose R into BCNF. Show keys, projected FDs.

Third Normal Form (3NF)



Let R be a relation schema, S be the set of FDs given to hold over R .

R is in 3NF if, for every FD $A_1A_2 \dots A_n \rightarrow B_1B_2 \dots B_m$

one of the following statements is true:

- same as
for BCNF* {
1. The FD is **trivial**: $\{B_1, B_2, \dots, B_m\} \subseteq \{A_1, A_2, \dots, A_n\}$
 2. $A_1A_2 \dots A_n$ is a candidate key of R
 3. $A_1A_2 \dots A_n$ is a superkey of R
 4. Each B_i is part of some candidate key of R

In contrast to BCNF, some redundancy is possible with 3NF. This normal form is a compromise, needed when no dependency-preserving decomposition into BCNF exists.

Decomposition into 3NF

Let R be a relation schema, \mathbf{S} be the set of FDs given to hold over R . We decompose R by considering FDs that violate 3NF.

Algorithm *3NFSynthesisDecomposition*

Input: Relation R , a set of FDs \mathbf{S} that hold in R .

Output: A decomposition of R into a set of relations, all of which are in 3NF.

1. Check whether R is in 3NF. If so, return $\{R\}$.
2. Find a minimal basis for \mathbf{S} , say \mathbf{T} .
3. For each FD in \mathbf{T} of the form $A_1A_2 \dots A_n \rightarrow B_1B_2 \dots B_m$ create a relation $A_1A_2 \dots A_nB_1B_2 \dots B_m$ and add it to the decomposition
4. If none of the relations from Step 3 is a key for R , another relation to the decomposition, whose schema is a key for R

Example

$$R(A, B, C, D, E) \quad S = \{AB \rightarrow C, C \rightarrow B, A \rightarrow D\}$$

Let's start by computing T , the minimum basis of S .

To do this, we compute attribute closures that appear in S , and the FDs to which they give rise.

$$\{A\}^+ = \{A, D\}^+ = \{A, D\}$$

$$A \rightarrow D$$

$$\{C\}^+ = \{B, C\}^+ = \{B, C\}$$

$$C \rightarrow B$$

$$\{B\}^+ = \{B\} \quad \{D\}^+ = \{D\} \quad \{E\}^+ = \{E\}$$

$$\{A, B\}^+ = \{A, C\}^+ = \{A, B, C, D\}$$

$$AB \rightarrow C \quad AB \rightarrow D \quad AC \rightarrow B \quad AC \rightarrow D$$

And now pick the minimal set of FDs. $T = \{AB \rightarrow C, A \rightarrow D, C \rightarrow B\}$

For example, we don't add $AB \rightarrow D$ because $A \rightarrow D$ is in T .

Example (continued 1)

$R(A, B, C, D, E)$

$S = \{AB \rightarrow C, C \rightarrow B, A \rightarrow D\}$

On the previous slide, we computed the minimum basis of S and called it T .

$T = \{AB \rightarrow C, A \rightarrow D, C \rightarrow B\}$

$\{A, B, E\}^+ = \{A, B, C, E, D\}$

Next, what are the candidate keys of R ?

$\{A, C, E\}^+ = \{A, B, C, E, D\}$

Is R in 3NF? No, the following FD violates 3NF: it's non-trivial, A is not a candidate key or a superkey, and D is not part of any candidate key.

$A \rightarrow D$

Decompose R into 3NF by synthesis.

First add a relation for each FD in T , shown here with the corresponding FDs.

$R_1(ABC) \quad AB \rightarrow C$

$R_2(AD) \quad A \rightarrow D$

$R_3(BC) \quad C \rightarrow B$

Observe that R_3 is a proper subset of R_1 . We never need to create a relation that is a proper subset of another, keep only R_1 and R_2 .

$R_1(ABC) \quad AB \rightarrow C, \quad C \rightarrow B$

$R_2(AD) \quad A \rightarrow D$

Example (continued 2)

$R(A, B, C, D, E)$

$S = \{AB \rightarrow C, C \rightarrow B, A \rightarrow D\}$

Two slides ago we computed the minimum basis of S , and called it T .

$T = \{AB \rightarrow C, A \rightarrow D, C \rightarrow B\}$

We then computed the candidate keys of R .

$\{A, B, E\}^+ = \{A, B, C, E, D\}$

$\{A, C, E\}^+ = \{A, B, C, E, D\}$

On the previous slide we computed the decomposition into 2 relations.

$R1(ABC) \quad AB \rightarrow C, \quad C \rightarrow B$

$R2(AD) \quad A \rightarrow D$

Finally, we need to add a relation that represents some (one) candidate key of R , since neither $R1$ nor $R2$ includes a candidate key of R . Include either $R3.1$ or $R3.2$ into the final decomposition.

$R3.1(ABE)$

$R3.2(ACE)$

Example

Are these relations in BCNF? In 3NF?

$R(ABCD) \ A \rightarrow B ; B \rightarrow A ; A \rightarrow D ; D \rightarrow B$

$R(ABCD) \ AB \rightarrow C ; BCD \rightarrow A ; D \rightarrow A ; B \rightarrow C$

$R(ABCD) \ FD's : AC \rightarrow D ; D \rightarrow A ; D \rightarrow C ; D \rightarrow B$

Example (solutions)

Are these relations in BCNF? In 3NF?

$R(ABCD) \ A \rightarrow B ; B \rightarrow A ; A \rightarrow D ; D \rightarrow B$

$R(ABCD) \ AB \rightarrow C ; BCD \rightarrow A ; D \rightarrow A ; B \rightarrow C$

$R(ABCD) \ FD's : AC \rightarrow D ; D \rightarrow A ; D \rightarrow C ; D \rightarrow B$

Schema refinement summary

- We discussed schema redundancy
- We presented functional dependencies (FDs), talked about how we can reason about them and use them to determine if a schema exhibits redundant
- Normal forms: BCNF and 3NF
- Properties of a decomposition: (1) eliminate redundancy; (2) lossless join; (3) dependency-preserving
- Algorithms for: closure of a set of attributes, projection of a set of FDs, BCNF decomposition, 3NF synthesis